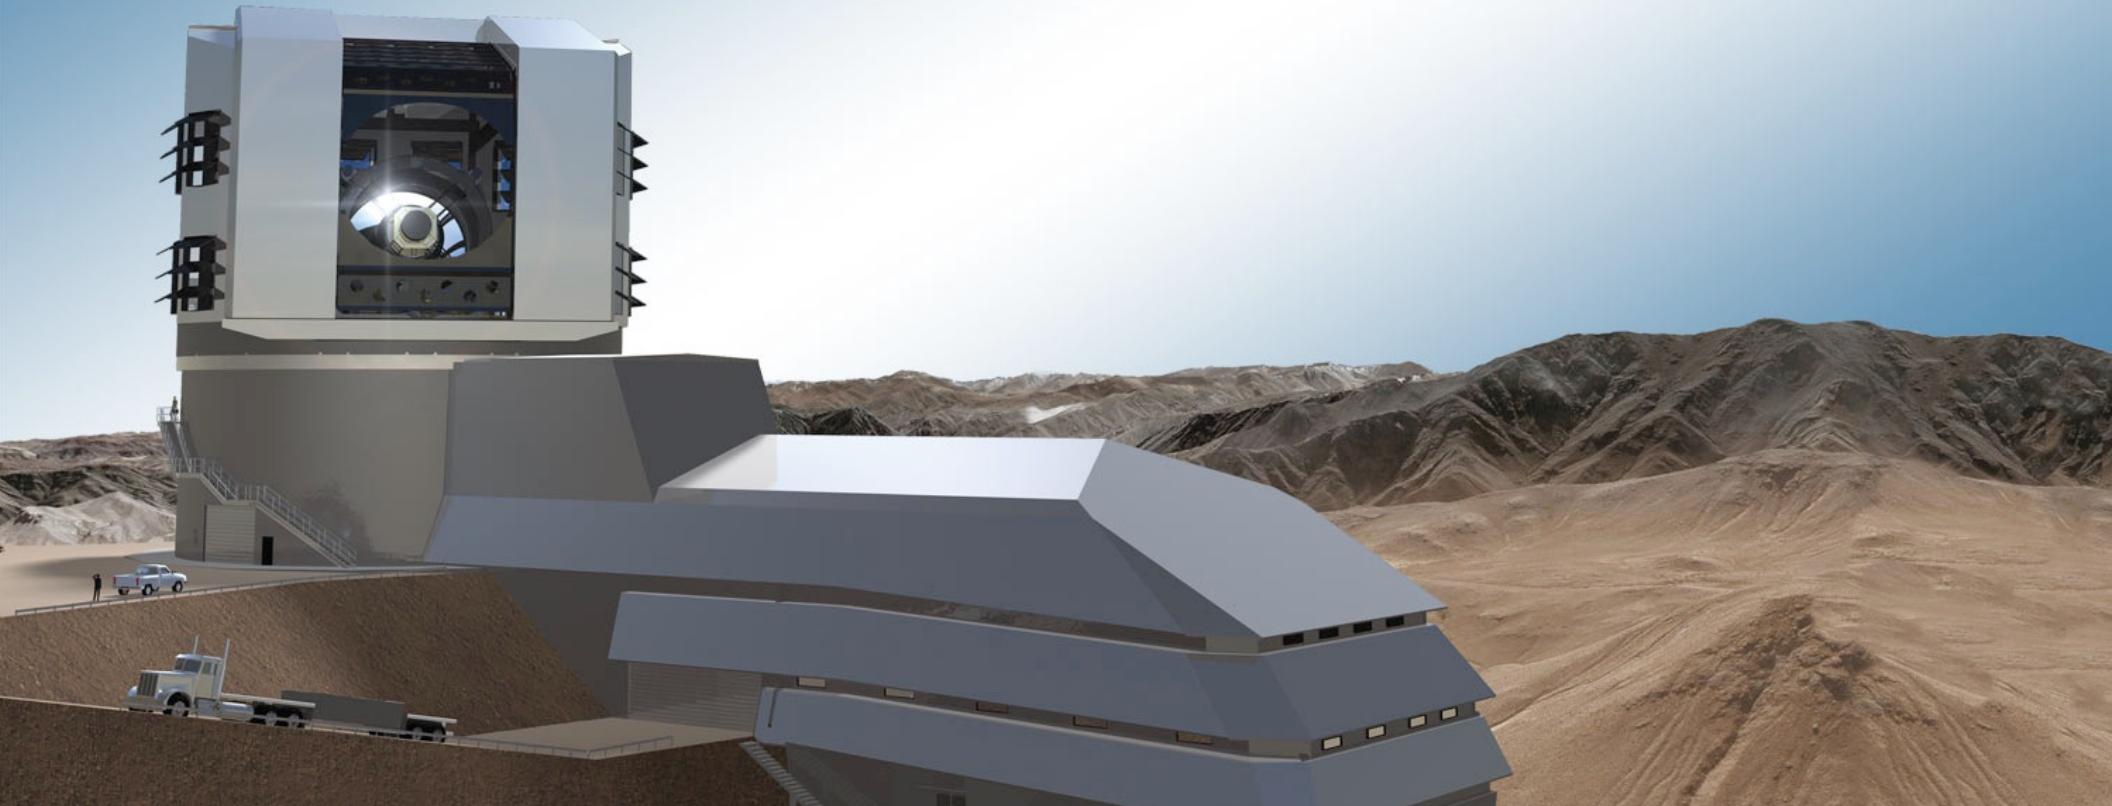




Photometric Redshifts for LSST

Jeffrey Newman, U. Pittsburgh / PITT-PACC

Deputy Spokesperson, LSST Dark Energy Science Collaboration



Before we start



- Please download *data_trim.fits.gz* and *Photo-z lessons.ipynb* from our Confluence site to a working directory, if you haven't done so already

Outline

- Overview of photometric redshifts
 - Template methods
 - Training-based methods
- Some open issues
 - Spectroscopic incompleteness
 - Robust training
 - $p(z)$ coverage
 - Combining results from multiple codes
 - $p(z,\alpha)$ storage
 - Optimizing spectroscopic samples
 - Defining ideal LSST algorithm
- Some examples of issues with current codes

LSST constrains dark energy in many ways... all will rely on redshift information

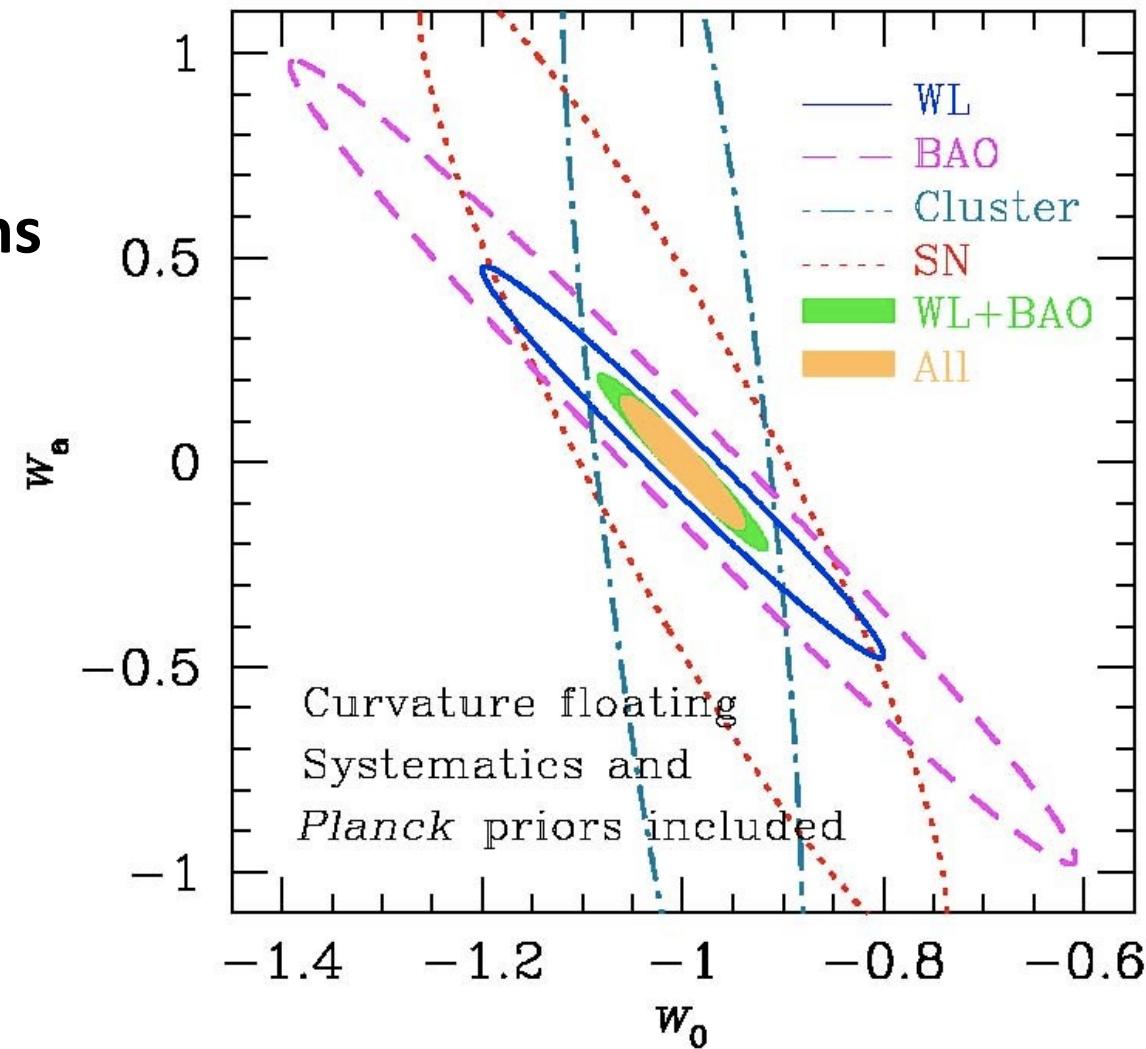


- LSST will constrain dark energy

via 4 major probes:

- Weak gravitational lensing
 - Baryon Acoustic Oscillations
 - Type Ia supernovae
 - Cluster counts
- (Plus strong lensing, etc.)

- For **all** of these, as well as
much galaxy and AGN science,
we want to measure some
observable as a function of
redshift



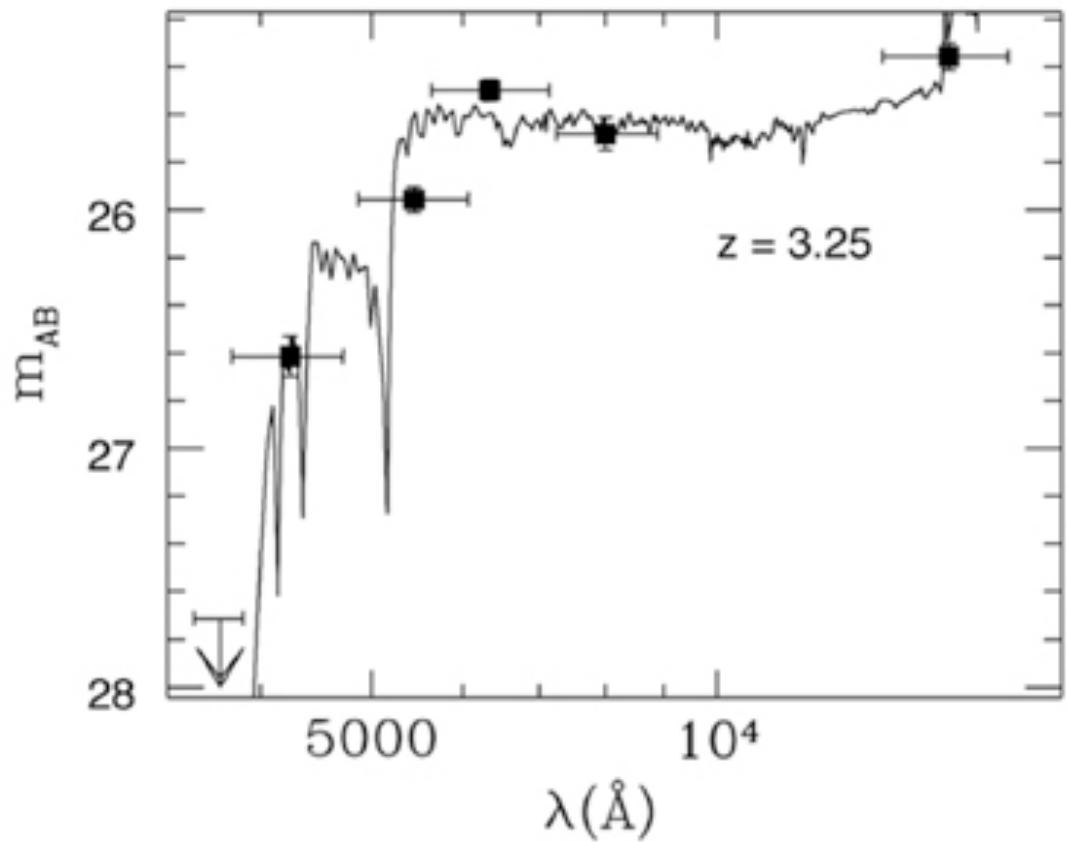
LSST Science Book

Spectroscopy provides ideal redshift measurements – but is infeasible for large samples

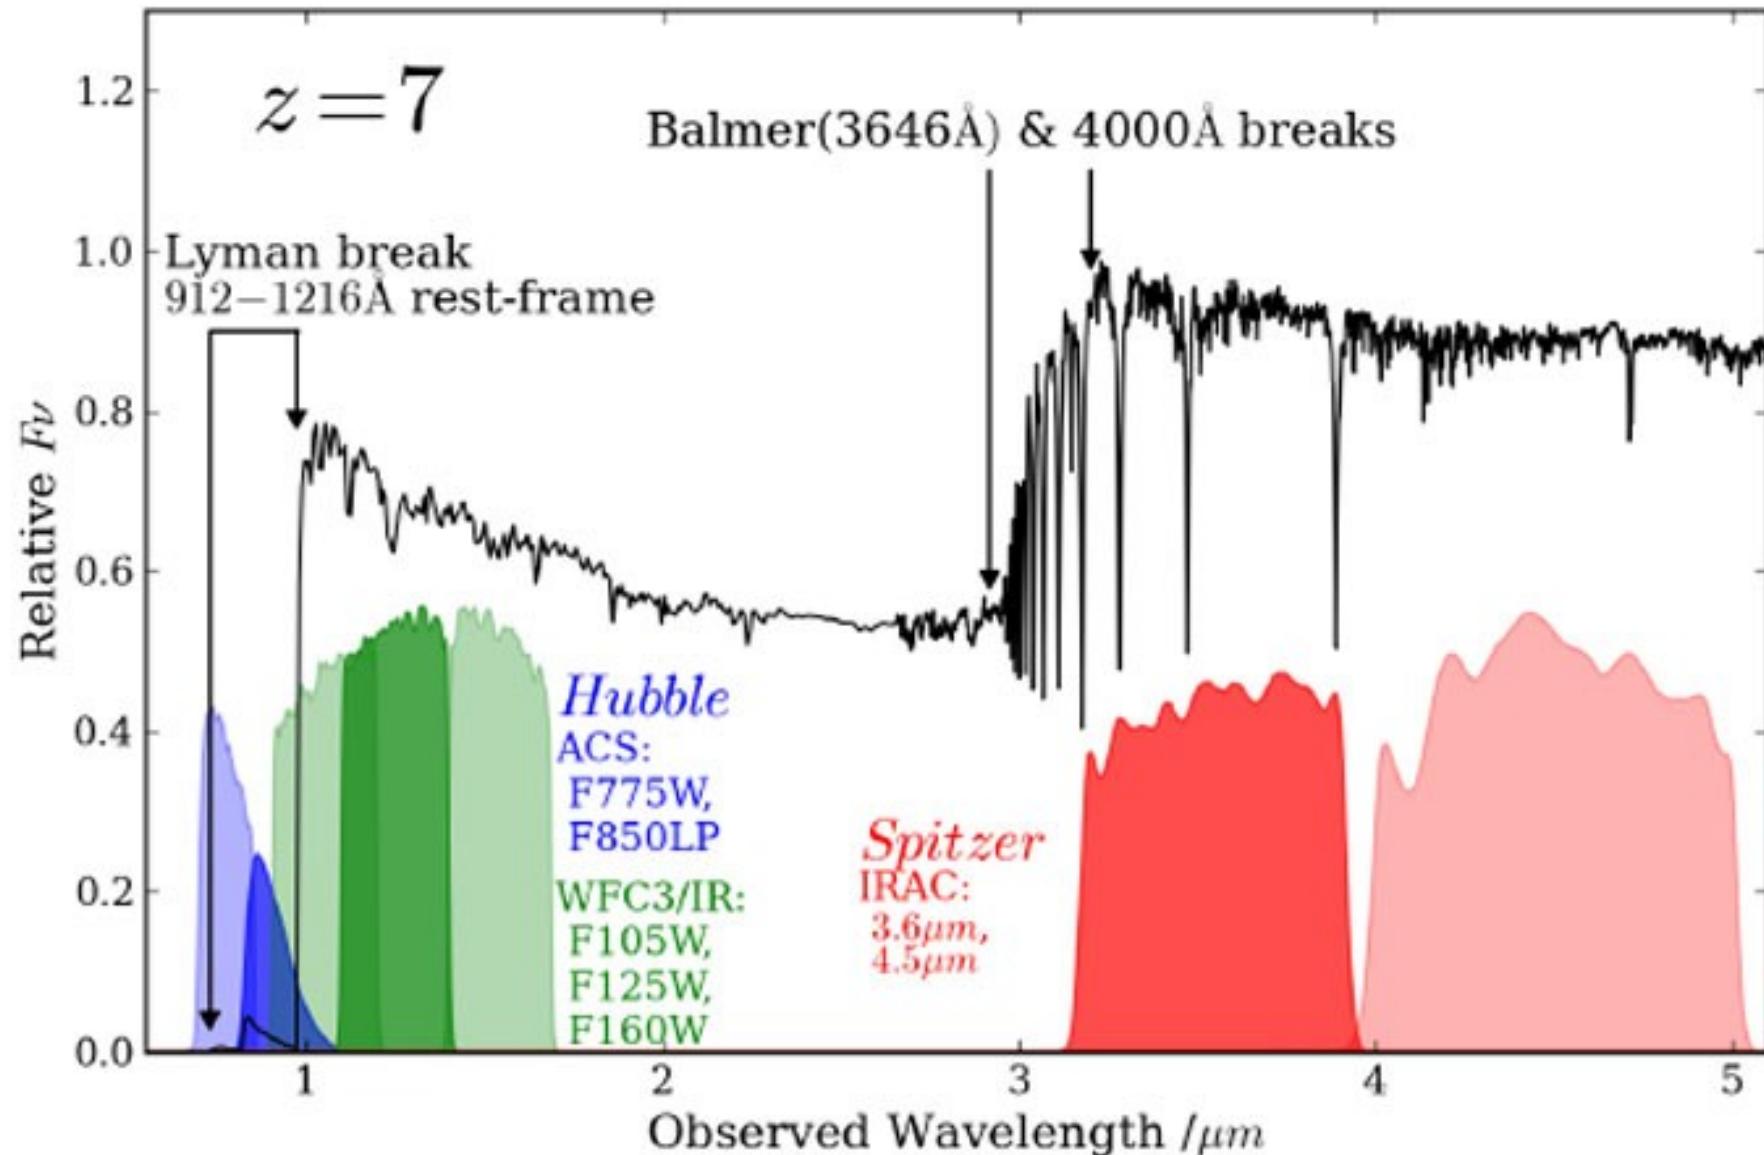
- Redshift ('z') measurements allow us to determine how far back in Universe's history we are looking for an object
- Study galaxy evolution, cosmology, etc. by measuring properties as a function of redshift
- To determine: measure spectrum of light from object with spectrograph; compare observed wavelengths of spectral features to rest frame values to get z
- At LSST “gold sample” ($i < 25.3$) depths, ~100 hours on a 10m telescope to determine a redshift (75% of time) spectroscopically
- With a next-generation, 5000-fiber spectrograph on a 10m telescope, still **>50,000 telescope-years** to measure redshifts for LSST “gold” weak lensing sample (4 billion galaxies)!

Spectroscopy provides ideal redshift measurements – but is infeasible for large samples

- Alternative: use broad spectral features to determine z : a **photometric redshift** or **photo-z**
- **Advantage:** high multiplexing
- **Disadvantages:** lower precision, calibration uncertainties



Photometric redshifts rely on the existence of broad spectral features in galaxy spectra...

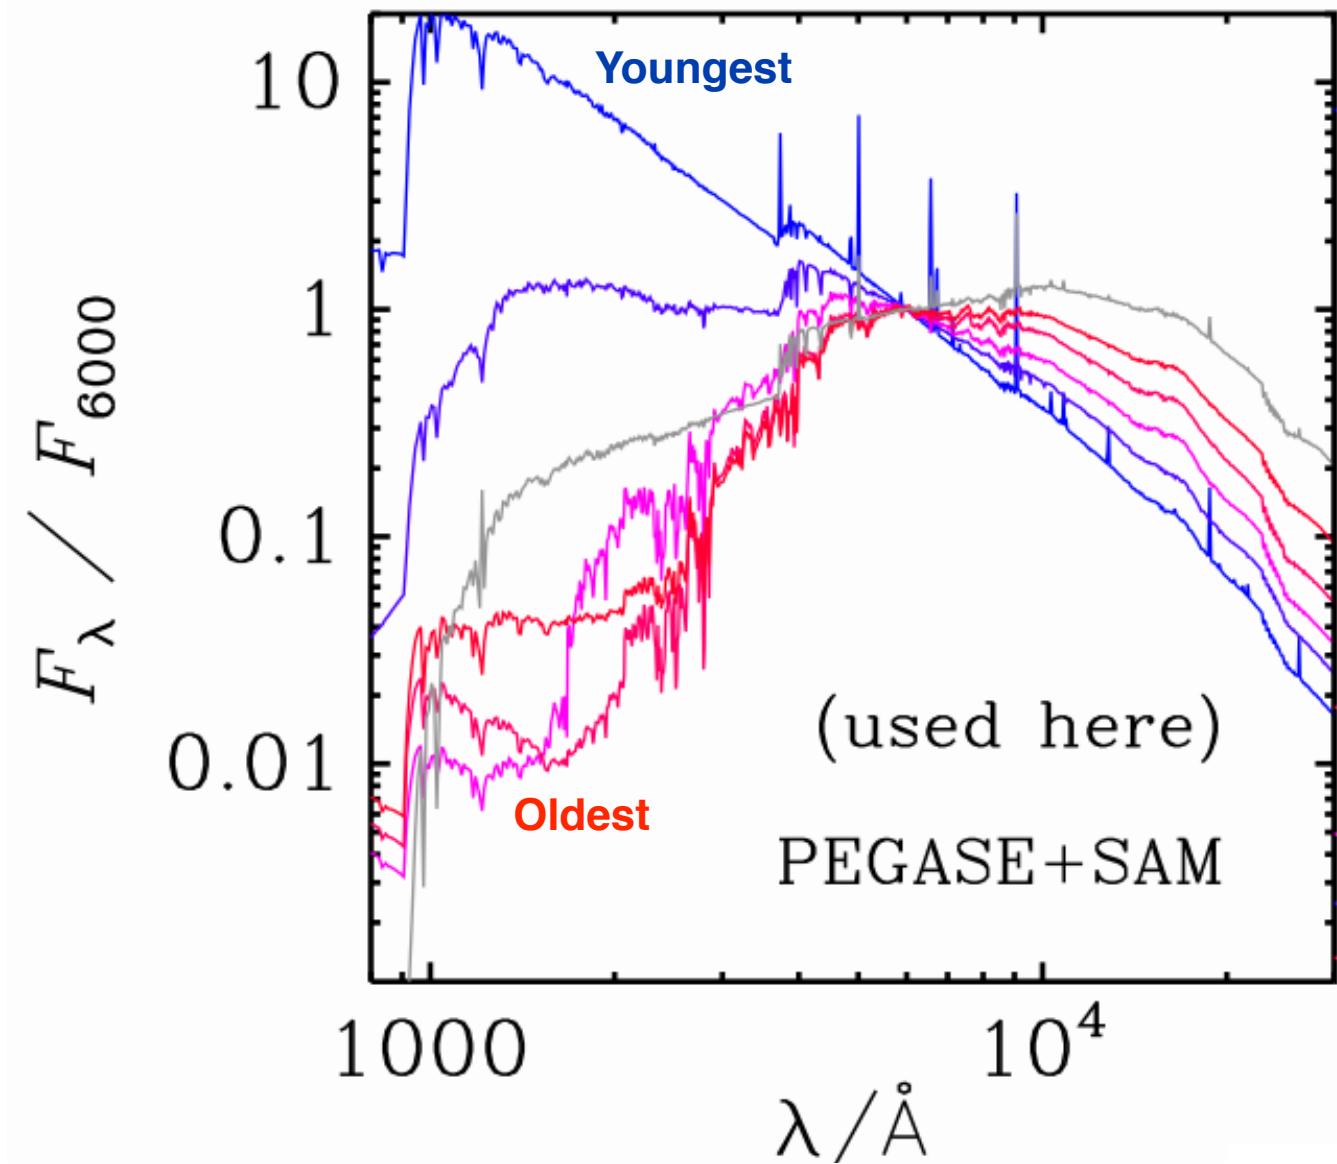


Dunlop 2012

but those features are stronger in some galaxies than others

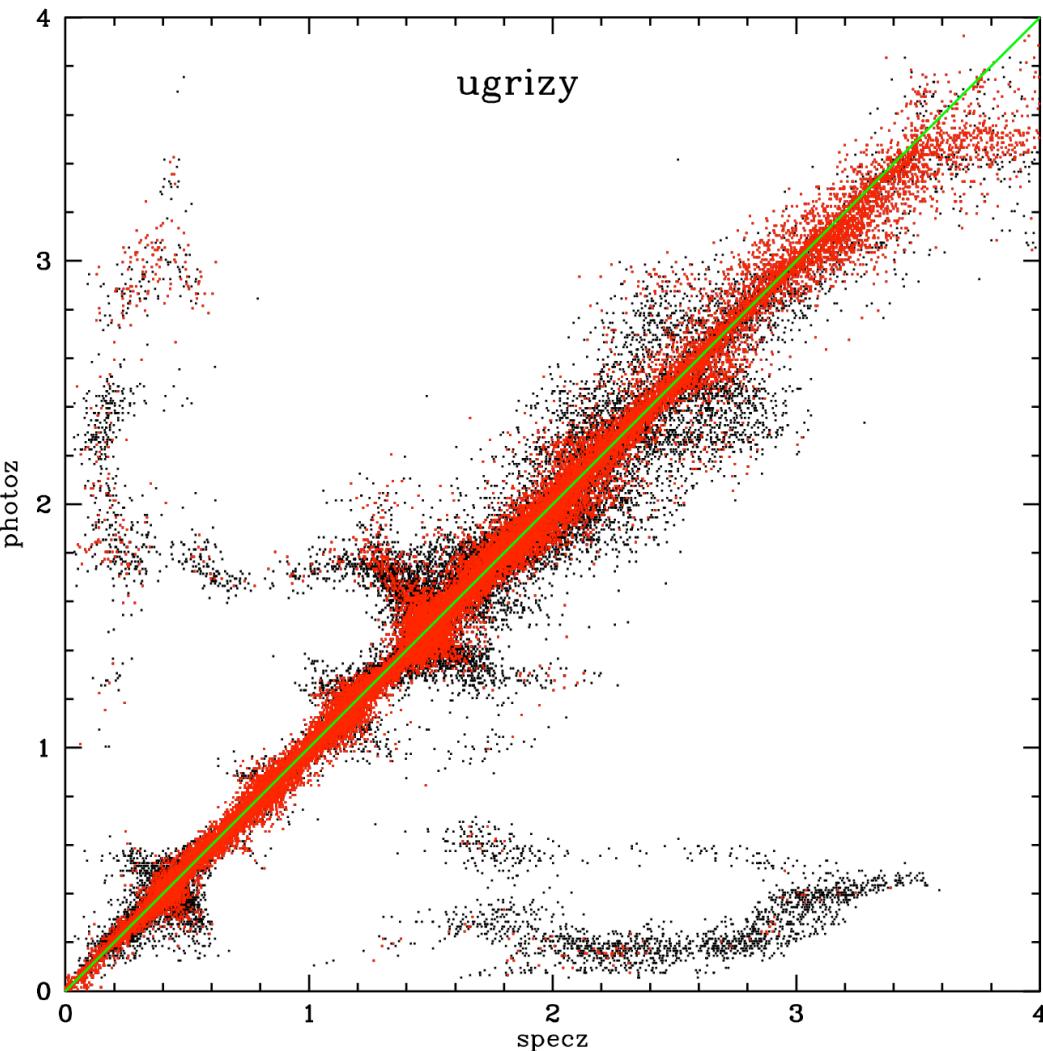


- Galaxies with older stellar populations exhibit stronger 'breaks'
- As a result, photo-z's can be more precise for redder galaxies
- At higher redshifts, blue galaxies with young stellar populations dominate - photo-z problem gets harder



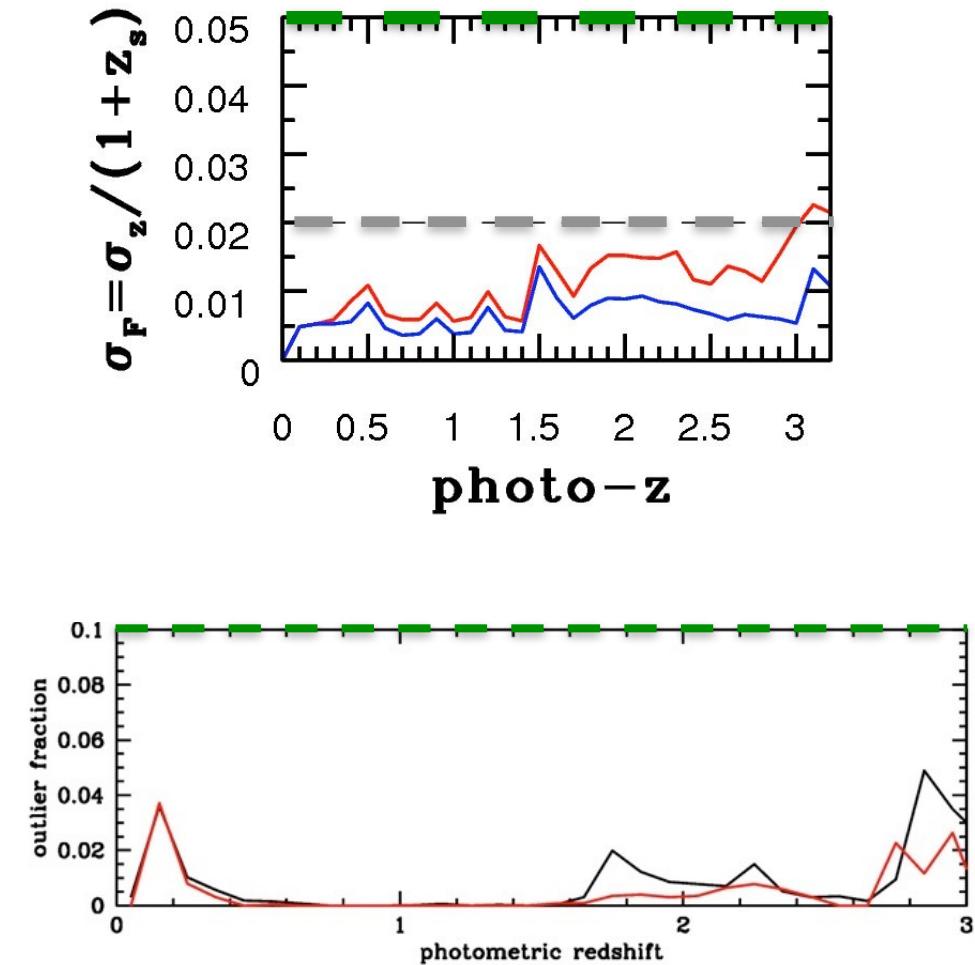
Brammer et al. 2008

Example: expected photo-z performance for LSST *ugrizy*



S. Schmidt

Green: Requirements on actual performance; grey: requirements on performance with perfect template knowledge (as in these sims)

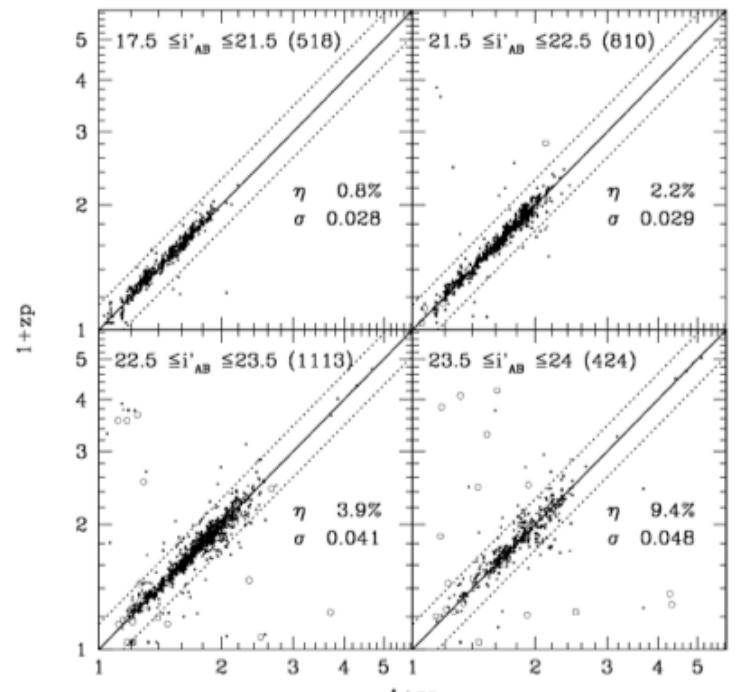


Basic methods: Template fitting photo-z's

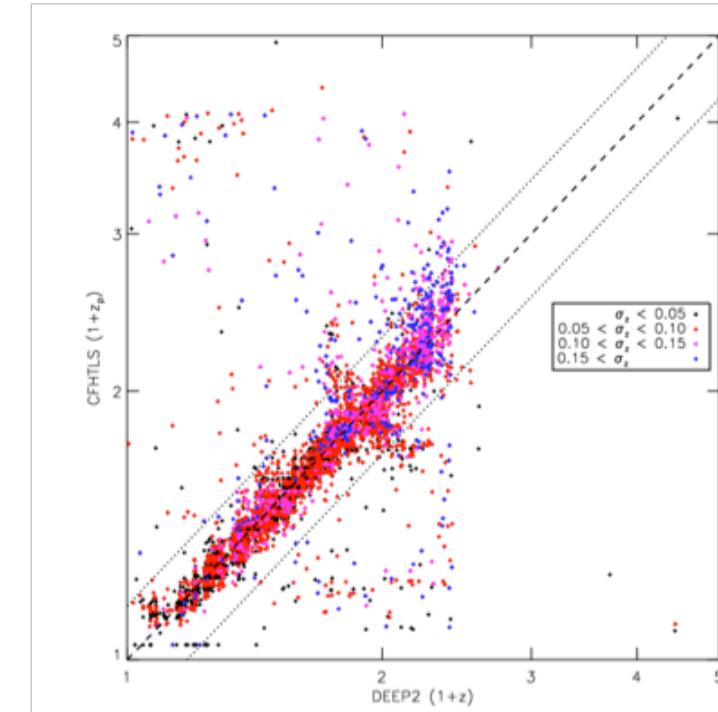


- Use galaxies with known z to calibrate set of underlying galaxy spectral energy distributions (SEDs) and photometric band-passes
 - Determine posterior probability distribution for $z \mid ugrizy$
 - Also provides info on galaxy properties from template fit

Needs spectra of galaxies spanning full range of possible properties to tune templates, establish priors, etc.



Ilbert et al. 2006



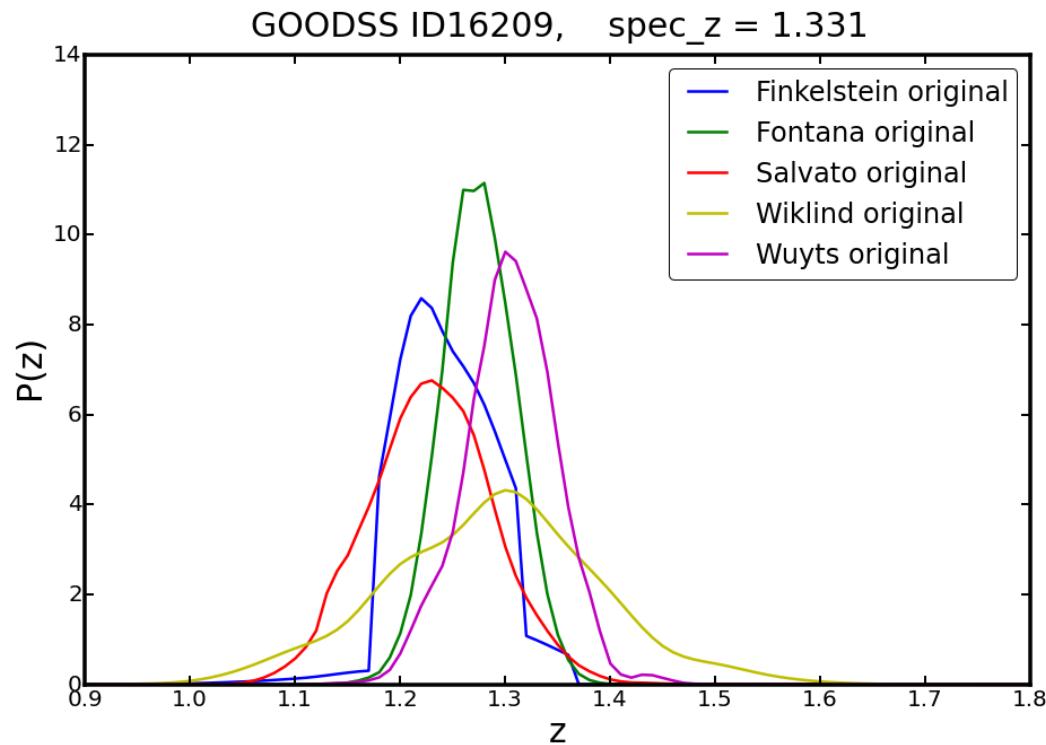
Ilbert photo-z's vs. DEEP2 z

Plot by Ben Weiner

Describing photometric redshift measurements



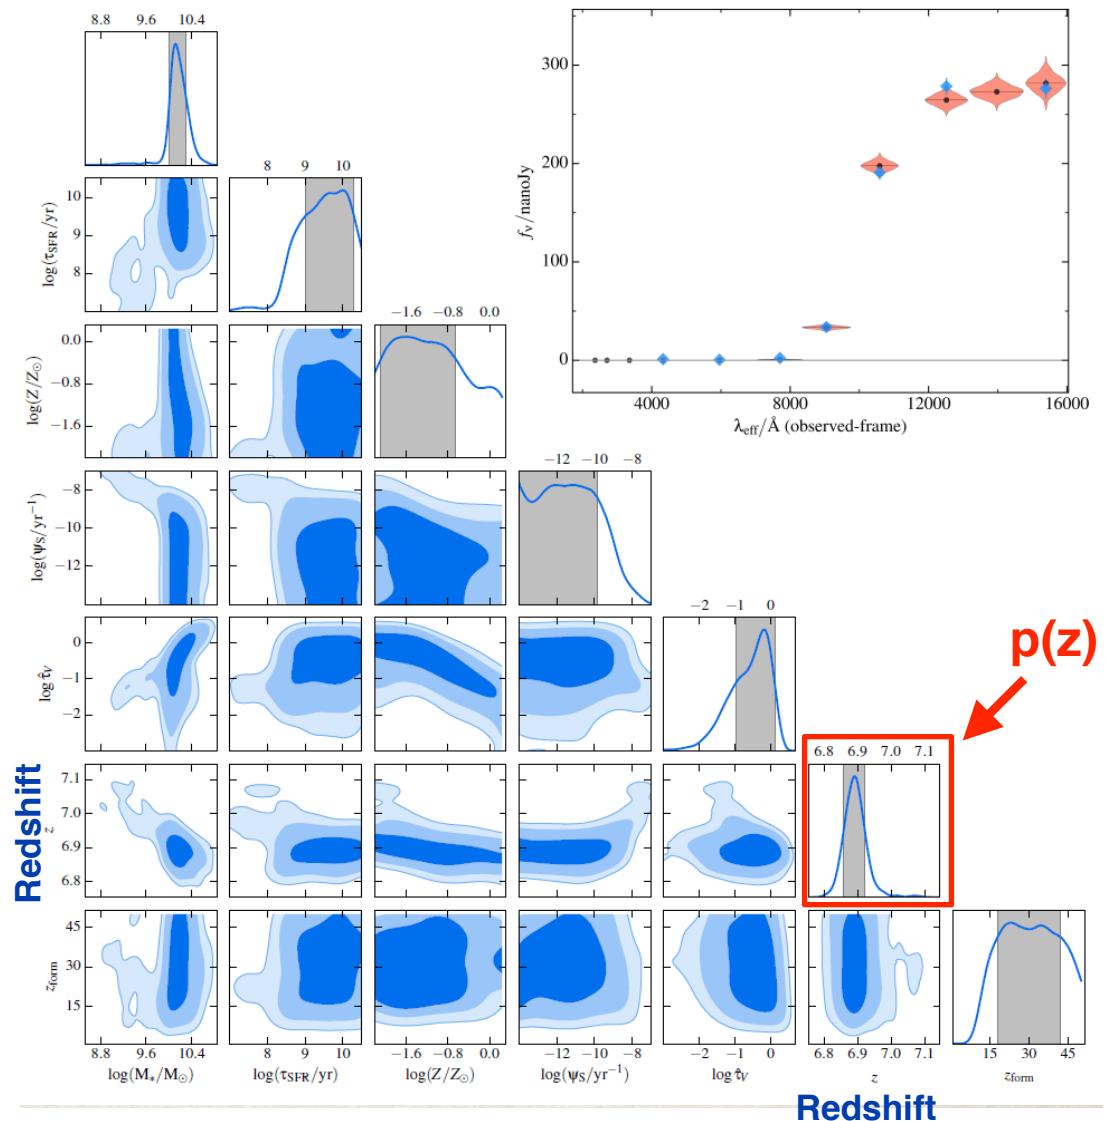
- Some codes simply output the best-fit z with errors
- Generally better to use the posterior probability distribution for $z \mid \text{fluxes}$: $p(z)$
- probability that $a < z < b$ =
 $\int_a^b p(z) dz$
 - $\int_0^\infty p(z) dz = 1$
- Various definitions for 'point' / single estimate from $p(z)$:
 - Peak of distribution
 - Expectation value of z (possibly only calculating using highest peak)



Describing photometric redshift measurements



- Can also provide info on galaxy properties from template fit
- E.g., template index T or galaxy parameters α_i such as stellar mass, star formation rate, etc.):
 $p(z, \alpha)$



Chevallard & Charlot 2016

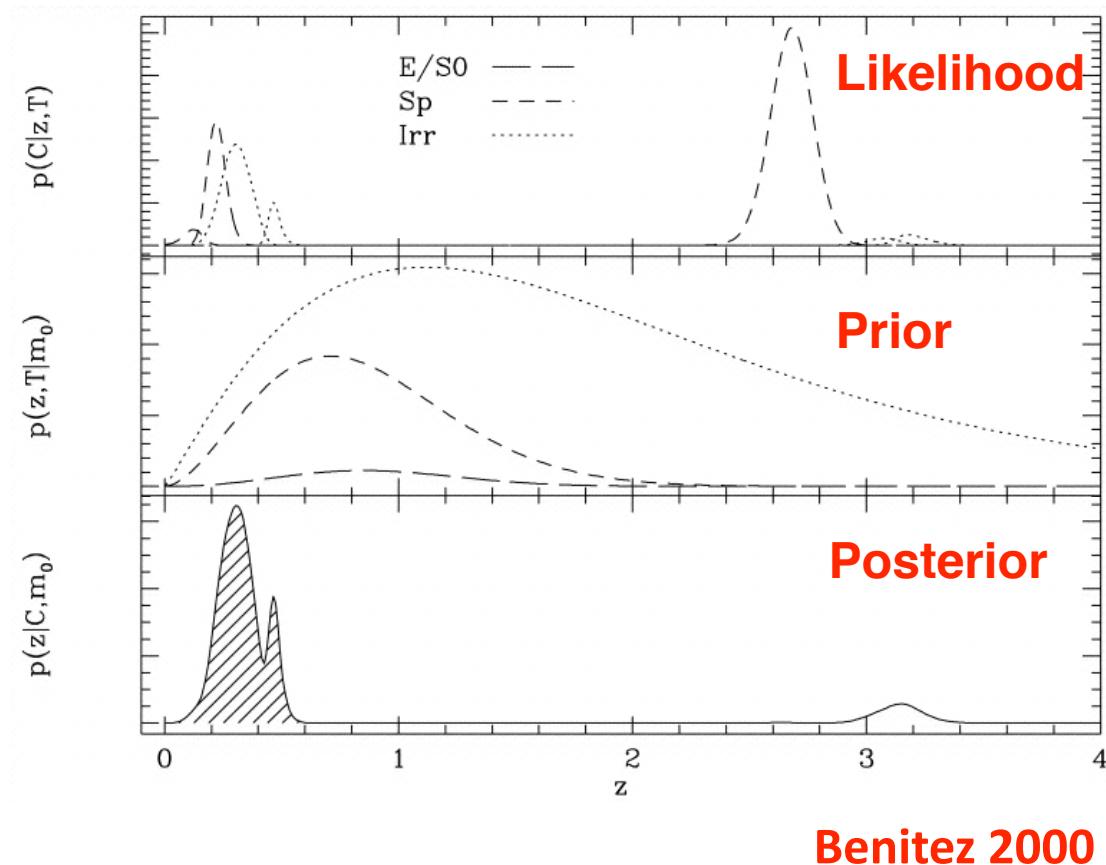
**Question to discuss in groups: What quantities do
you want for each LSST galaxy?**



Basic methods: Template fitting photo-z's



- Typical algorithms:
 - Determine **likelihood** of colors (=ratios of fluxes between bands) as a function of z and template
 - Often via $\chi^2(z, T)$ or $\min(\{\chi^2(z | T)\})$; some algorithms use linear combinations of templates
 - Typically utilize **prior** for redshift or redshift & type based on magnitude (sometimes size/morphology as well)
 - Then multiply to get **posterior**...



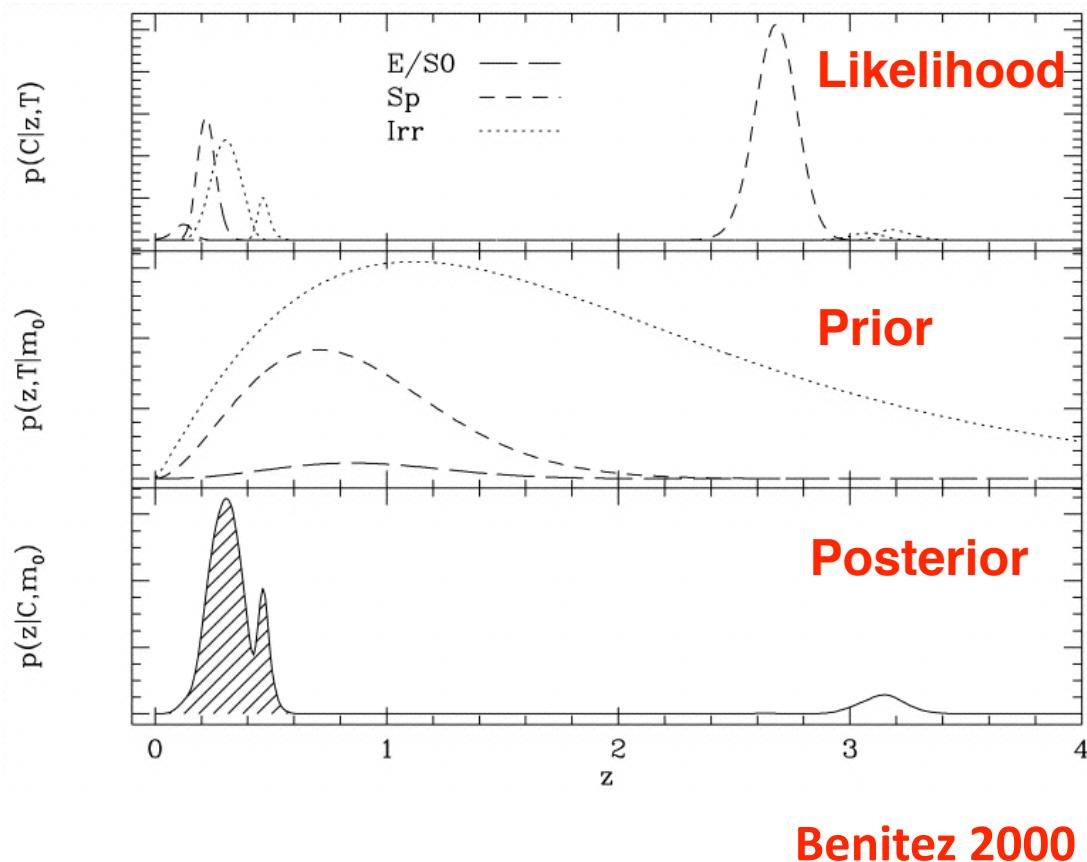
Benitez 2000

Use spectra of galaxies spanning full range of possible properties to tune templates/filter systems, establish priors, etc.

Basic methods: Template fitting photo-z's



- Note that this is standard Bayesian language:
- $p(\text{fluxes} | z) = \text{likelihood}$
- $p(z) = \text{prior}$
- $p(z | \text{fluxes}) = \text{posterior}$
- By Bayes' theorem,
$$p(z | \text{fluxes}) = \frac{p(\text{fluxes} | z) p(z)}{p(\text{fluxes})}$$
- We often just normalize the integral of the posterior to be 1 rather than calculating $p(\text{fluxes})$

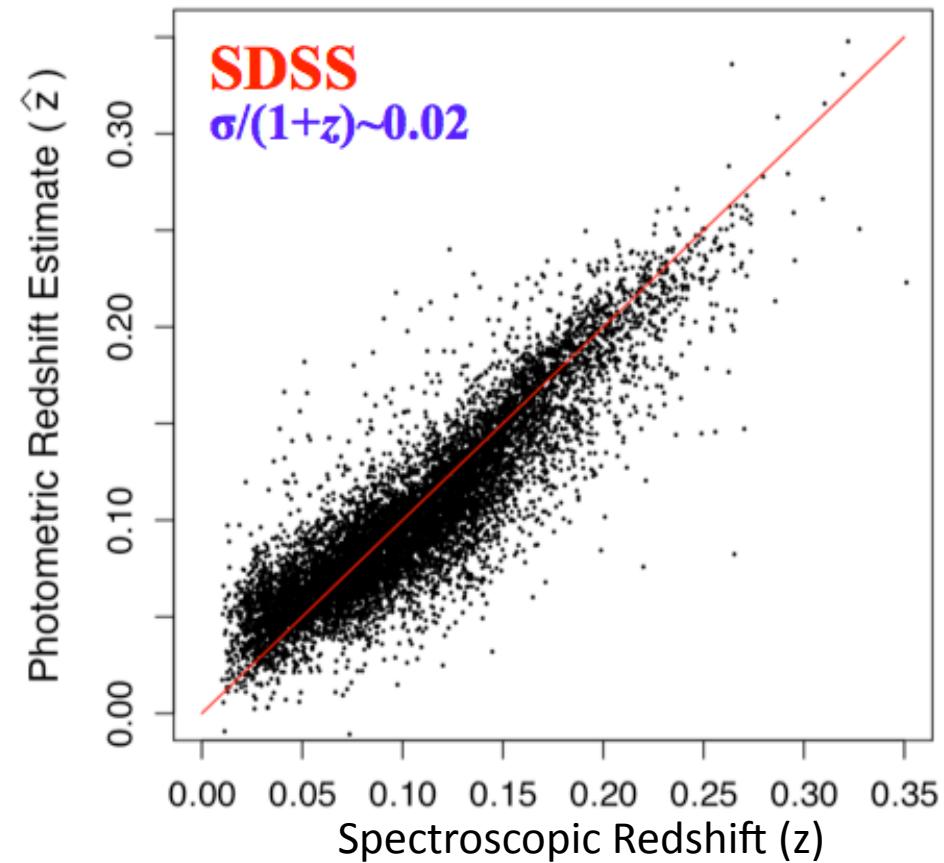


Benitez 2000

Basic methods: Training-based photo-z's



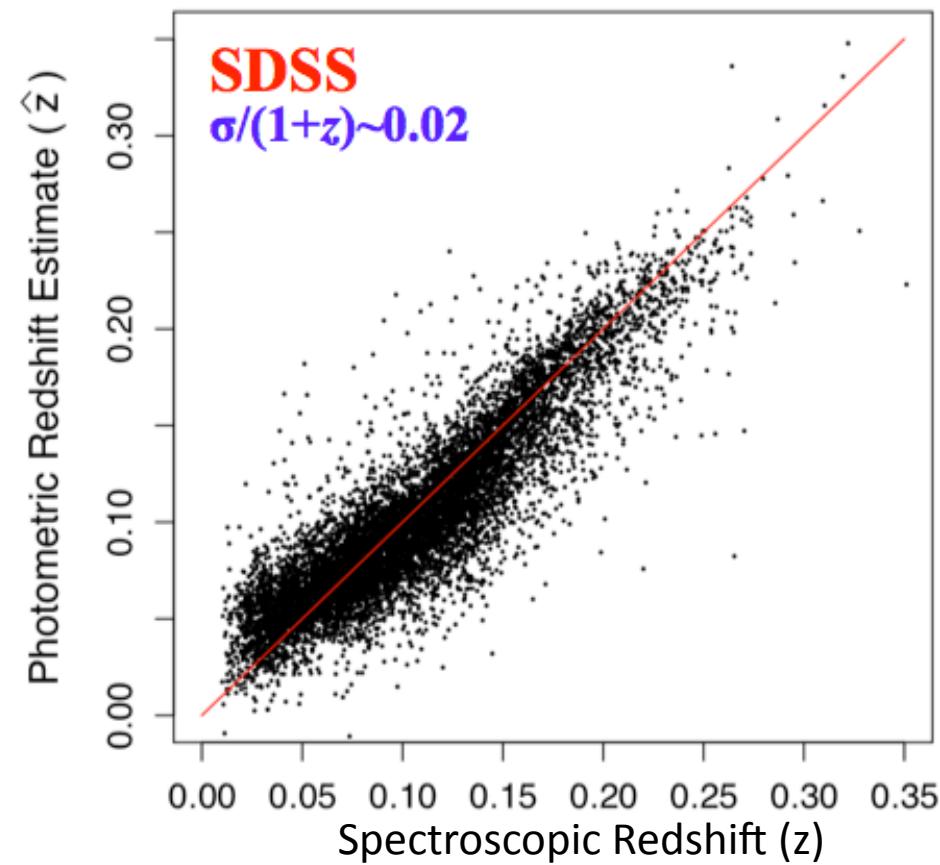
- Use galaxies with known redshift **and** uniform/well-understood sampling to determine relationship between z and colors/fluxes
- Can take advantage of progress in machine learning & stats, but generally **extrapolate poorly**; Training set **MUST** span full range of properties & z of galaxies



Basic methods: Training-based photo-z's



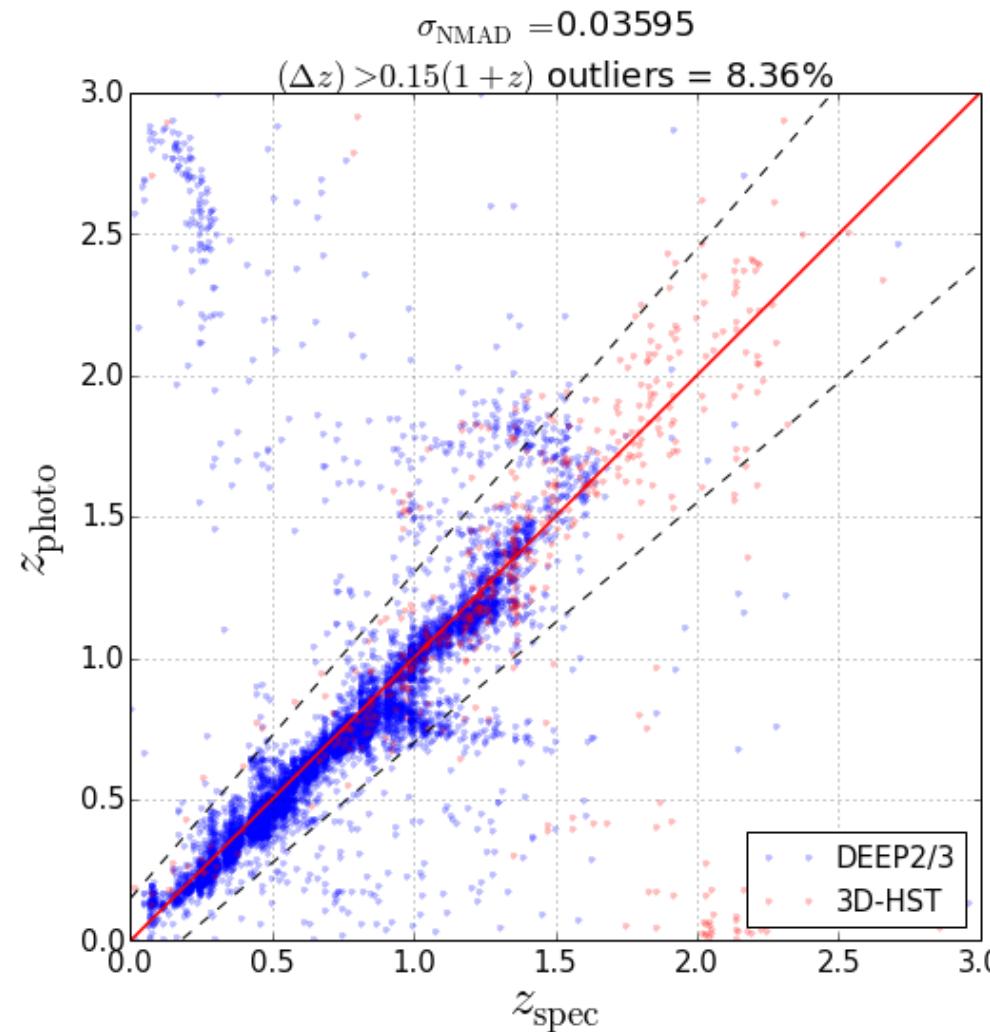
- Many algorithms: e.g.
 - Neural networks
 - Boosted Decision Trees
 - Random Forest regression
 - k-Nearest Neighbor
 - Diffusion map + regression
- For bright, nearby galaxies, training sets are ~complete and both template-based & training-set-based algorithms perform extremely similarly



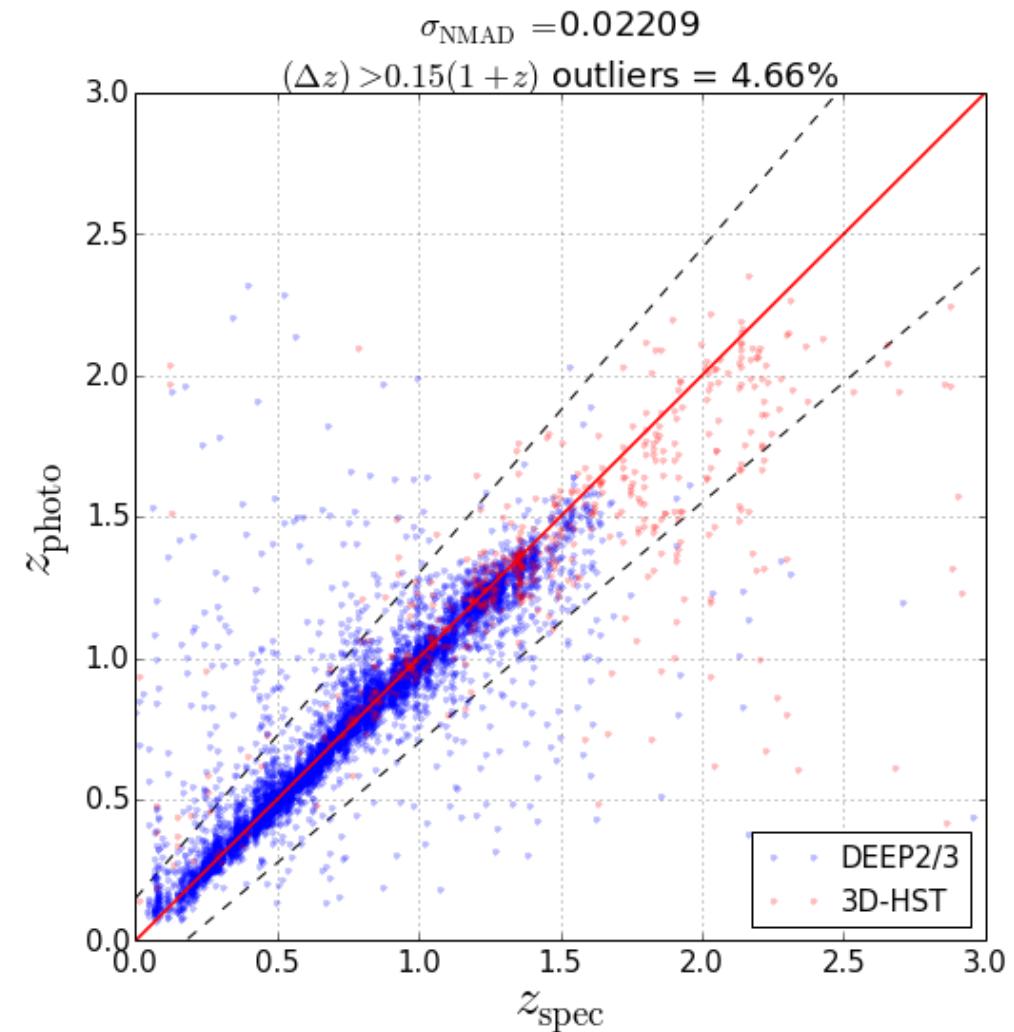
At higher redshifts, the photo-z problem is more difficult

- Zhou et al. 2016 (in prep.): empirical, LSST-like dataset: CFHT LS *ugriz* + Subaru *y* + DEEP2/DEEP3/3D-HST redshifts

EAZY (template code, untuned)



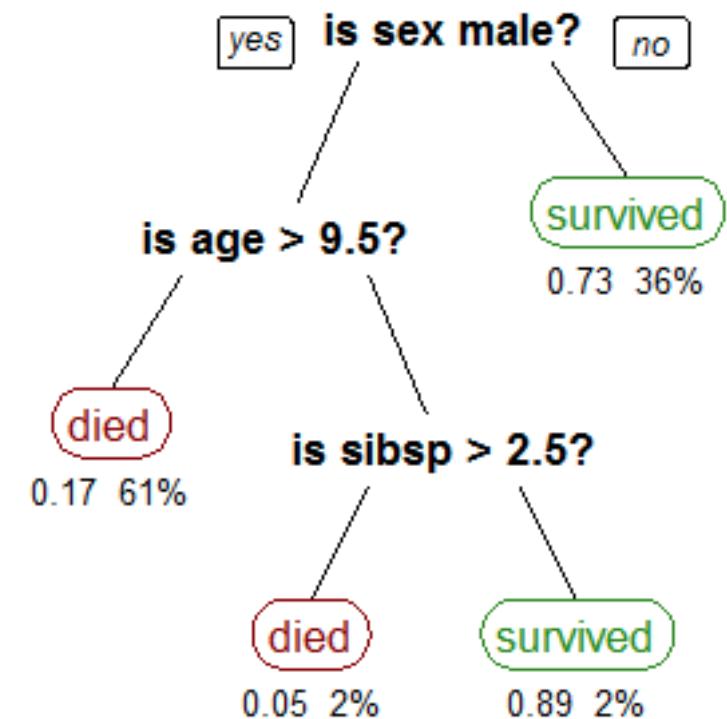
Random Forest Regression



Let's try determining our own photo-z's!



- We'll apply the random forest regression routine in the scikit-learn Python package to the dataset produced by Rongpu Zhou.
- Random forest is a highly effective machine learning technique based on decision trees
- In random forest, a wide variety of decision trees get trained by a training set of data, randomly bootstrapping amongst that data for optimizing each tree
- A random subset of possible data features are used in each tree
- In the end, results from all trees get combined to give a prediction (in our case, a prediction for z)



Walking through the code: importing routines



- Load the Jupyter notebook (e.g. go to a Terminal; type 'jupyter notebook Photo-z\ Lessons.ipynb', or just type 'jupyter notebook' and choose the right file.)
- This notebook contains text and python code. To execute a block of code, click in the box of code, hold down shift, and press enter.
- We start by importing python modules and routines

```
[1]: %matplotlib inline
from __future__ import division #avoids integer division; 1/2 = 0.5, not 0
import numpy as np #numpy routines
from astropy.table import Table #astropy routine for reading tables
import matplotlib.pyplot as plt #plotting routines

# Random forest routine from scikit-learn:
from sklearn.ensemble import RandomForestRegressor

# Cross-Validation routines:
from sklearn.cross_validation import KFold
from sklearn.cross_validation import train_test_split
from sklearn.cross_validation import cross_val_predict
```

Walking through the code: reading in the data with Table.read()



```
In [2]: #read in catalog of magnitudes and redshifts
cat = Table.read('data_trim.fits.gz')

# How big is the catalog?
print 'Full catalog: ',len(cat), ' objects'

# What information does it contain for each object?
print cat.colnames
print cat
```

```
Full catalog: 8508 objects
['U', 'G', 'R', 'I', 'Z', 'Y', 'UERR', 'GERR', 'RERR', 'IERR', 'ZERR',
 'YERR', 'RADIUS_ARCSEC', 'ZHELIO', 'MAGB', 'UB_0']
```

- **U, G, R, I, Z, and Y** are magnitude measurements for each object (=2.5 times the \log_{10} of the flux relative to some reference flux, the "zero point"). **Is everyone familiar with the idea of magnitudes?**
- **UERR, GERR, RERR, IERR, ZERR, and YERR** are uncertainties in the magnitudes

Walking through the code: reading in the data with Table.read()



```
In [2]: #read in catalog of magnitudes and redshifts
cat = Table.read('data_trim.fits.gz')

# How big is the catalog?
print 'Full catalog: ',len(cat), ' objects'

# What information does it contain for each object?
print cat.colnames
print cat
```

```
Full catalog: 8508 objects
['U', 'G', 'R', 'I', 'Z', 'Y', 'UERR', 'GERR', 'RERR', 'IERR', 'ZERR',
 'YERR', 'RADIUS_ARCSEC', 'ZHELIO', 'MAGB', 'UB_0']
```

- **RADIUS_ARCSEC** is a measure of the radius of the object in arc seconds
- **ZHELIO** is the redshift ($z = \lambda_{\text{observed}} / \lambda_{\text{restframe}} - 1$) in the heliocentric reference frame (why does that make a difference?)
- **UB_0** is an estimate of the restframe U-B color of the object (= -2.5 times the \log_{10} of the flux ratio between the U and B bands; bigger = redder)

Walking through the code: a function to make plots and calculate statistics



```
In [3]: #A function that we will call a lot: makes the zphot/zspec plot and calculates key statistics
def plot_and_stats(zspec,zphot):

    x = np.arange(0,5.4,0.05)
    outlier_upper = x + 0.15*(1+x)
    outlier_lower = x - 0.15*(1+x)

    mask = np.abs((z_phot - z_spec)/(1 + z_spec)) > 0.15
    notmask = ~mask

        #Standard Deviation of the predicted redshifts compared to the data:
    std_result = np.std((z_phot - z_spec)/(1 + z_spec), ddof=1)
    print 'Standard Deviation: %6.4f' % std_result

        #Normalized MAD (Median Absolute Deviation):
    nmad = 1.48 * np.median(np.abs((z_phot - z_spec)/(1 + z_spec)))
    print 'Normalized MAD: %6.4f' % nmad

        #Percentage of delta-z > 0.15(1+z) outliers:
    eta = np.sum(np.abs((z_phot - z_spec)/(1 + z_spec)) > 0.15)/len(z_spec)
    print 'Delta z >0.15(1+z) outliers: %6.3f percent' % (100.*eta)

        #Median offset (normalized by (1+z); i.e., bias:
    bias = np.median(((z_phot - z_spec)/(1 + z_spec)))
    sigbias=std_result/np.sqrt(0.64*len(z_phot))
    print 'Median offset: %6.3f +/- %6.3f' % (bias,sigbias)
```

Walking through the code: a function to make plots and calculate statistics



```
# make photo-z/spec-z plot
plt.figure(figsize=(8, 8))

#add lines to indicate outliers
plt.plot(x, outlier_upper, 'k--')
plt.plot(x, outlier_lower, 'k--')
plt.plot(z_spec[mask], z_phot[mask], 'r.', markersize=6, alpha=0.5)
plt.plot(z_spec[notmask], z_phot[notmask], 'b.', markersize=6, alpha=0.5)
plt.plot(x, x, linewidth=1.5, color = 'red')
plt.title('$\sigma_{\mathrm{NMAD}} = $\b{6.4f}\n'%'$ (\Delta z) > 0.15(1+z) $ outliers = %6.3f' %(eta*100) + ' ', fontsize=18)
plt.xlim([0.0, 2])
plt.ylim([0.0, 2])
plt.xlabel('$z_{\mathrm{spec}}$', fontsize = 27)
plt.ylabel('$z_{\mathrm{photo}}$', fontsize = 27)
plt.grid(alpha = 0.8)
plt.tick_params(labelsize=15)
plt.show()
```

Walking through the code: defining convenience arrays



```
# create convenience arrays for all magnitudes
u_mag = cat['U']
g_mag = cat['G']
r_mag = cat['R']
i_mag = cat['I']
z_mag = cat['Z']
y_mag = cat['Y']

# and for magnitude errors
u_err = cat['UERR']
g_err = cat['GERR']
r_err = cat['RERR']
i_err = cat['IERR']
z_err = cat['ZERR']
y_err = cat['YERR']

#Redshift array
z = cat['ZHELIO']

#Size array: radius in arcseconds
rad = cat['RADIUS_ARCSEC']

#Restframe color
ub_color = cat['UB_0']

#Photometry perturbed: doubling sizes of all errors
u_magn=u_mag + np.sqrt(3)*u_err*np.random.randn(len(u_mag))
g_magn=g_mag + np.sqrt(3)*g_err*np.random.randn(len(g_mag))
r_magn=r_mag + np.sqrt(3)*r_err*np.random.randn(len(r_mag))
i_magn=i_mag + np.sqrt(3)*i_err*np.random.randn(len(i_mag))
z_magn=z_mag + np.sqrt(3)*z_err*np.random.randn(len(z_mag))
y_magn=y_mag + np.sqrt(3)*y_err*np.random.randn(len(y_mag))
```

Walking through the code: putting data in the right format for scikit_learn regression routines



```
# Now, set up input data arrays for random forest regression

# First: magnitudes only
data_mags = np.column_stack((u_mag,g_mag,r_mag,i_mag,z_mag,y_mag))

# Next: colors only
data_colors = np.column_stack((u_mag-g_mag, g_mag-r_mag, r_mag-i_mag, i_mag-z_mag, z_mag-y_mag))

# Next: colors and one magnitude
data_colmag = np.column_stack((u_mag-g_mag, g_mag-r_mag, r_mag-i_mag, i_mag-z_mag, z_mag-y_mag, i_mag))
perturbed_colmag=np.column_stack((u_magn-g_magn, g_magn-r_magn, r_magn-i_magn, i_magn-z_magn, z_magn-y_magn, i_magn))

# Finally: colors, magnitude, and size
data_colmagsize = np.column_stack((u_mag-g_mag, g_mag-r_mag, r_mag-i_mag, i_mag-z_mag, z_mag-y_mag, i_mag, rad))

data_z = z

# We need to set up an implementation of the scikit-learn RandomForestRegressor in an object called 'regrn'.
regrn = RandomForestRegressor(n_estimators = 50, max_depth = 30, max_features = 'auto')
```

- Not particularly interesting... scikit-learn routines expect the features to be used for predictions to be in a 2d array, with dimensions (# of objects, # of features)

Next: plotting colors, magnitude and size as a function of redshift



- Note: to an astronomer, color is the difference between magnitudes in 2 filters (with bluest specified first); e.g. $g - r = g$ -band magnitude - r -band magnitude = $-2.5 \log_{10} (g \text{ flux} / r \text{ flux})$
- Color is larger for redder objects
- Some photo-z codes work in flux (better for handling uncertainties, which are closer to Gaussian in flux)

Next: plotting colors, magnitude and size as a function of redshift



```
: if 0:  
# Plots of colors vs. redshift  
  
plt.figure(figsize=(8, 5))  
  
plt.plot(z, u_mag-g_mag, '.', color = 'red', markersize=3, alpha=0.2)  
# Axis limits:  
plt.xlim([0., 1.5])  
plt.ylim([-0.5, 2.5])  
#Axis labels  
plt.ylabel('$u-g$ color', fontsize=20)  
plt.xlabel('Redshift', fontsize=20)  
#Add grid lines  
plt.grid()
```

- Change 'if 0' to 'if 1'. **What did we just do?**
- Take a look at all of the plots of color etc. vs. redshift. **Discuss in your groups:** what colors constrain redshift best for objects at $z \sim 0.2$? 0.6 ? 1 ? Why should different colors be important at different redshifts?

Next: plotting colors, magnitude and size as a function of redshift



```
if 0:
```

```
# To better assess the quality of the Random Forest fitting,  
# we split the data into Training (50%) and Test (50%) sets.  
# The code below performs this task on the data_mags and data_z arrays:  
data_train, data_test, z_train, z_test = train_test_split(data_mags, data_z,  
                                         test_size = 0.50, train_size = 0.50, random_state=7182016)  
  
#Train the regressor using the training data  
regrn.fit(data_train,z_train)  
  
#Apply the regressor to predict values for the test data  
z_phot = regrn.predict(data_test)  
z_spec = z_test  
  
#Make a photo-z/spec-z plot and output summary statistics.  
plot_and_stats(z_spec,z_phot)
```

- Change 'if 0' to 'if 1' .
- This does all the work: divides data into training and test sets, trains the random forest regression, uses it to predict redshifts for the test set, and has plots_and_stats do all the calculations.

Worse and better ways to do this



- **Worse: use the same data for training & test. Try it; discuss in your groups what went wrong.**
- **Better: use k-fold cross-validation. Compare the plot for 5-fold cross-validation to our first plot. What is different?**

In k-fold cross-validation, we split the data into k subsets. We loop over the subsets, training with all but one and testing with the other; in the end, we get the performance of training with a fraction $\frac{k-1}{k}$ of the data, but are able to get test statistics based on the *entire* dataset.

This is easy to do in scikit-learn, but does mean running the training k times on more data than before...

Exploiting the underlying structure of the data can improve machine learning results



- Next, we explore using different features to do the photo-z's:
 - Colors, instead of magnitudes
 - Add a magnitude (emulating effect of a prior)
 - Add size information
 - Before running the new ways, predict what will happen and discuss with your group.
 - How significant are the improvements?

Exploiting the underlying structure of the data can improve machine learning results



- Finally, let's explore doing photo-z's restricted to one particular population of galaxies
 - Red galaxies have strong breaks in their spectra and can yield tighter photo-z's ($\sigma \sim 0.01$).
 - Let's select the red population in our catalog. First, we plot the distribution of rest frame colors of galaxies:

```
if 0:  
    plt.figure(figsize=(8, 5))  
  
    inrange=(ub_color > -0.5) & (ub_color < 2)  
    plt.hist(ub_color[inrange],bins=100)  
    # Axis limits:  
    plt.xlim([-0.3, 1.5])  
    plt.xlabel('Restframe $U-B$ color', fontsize=20)  
    plt.grid()
```

Exploiting the underlying structure of the data can improve machine learning results

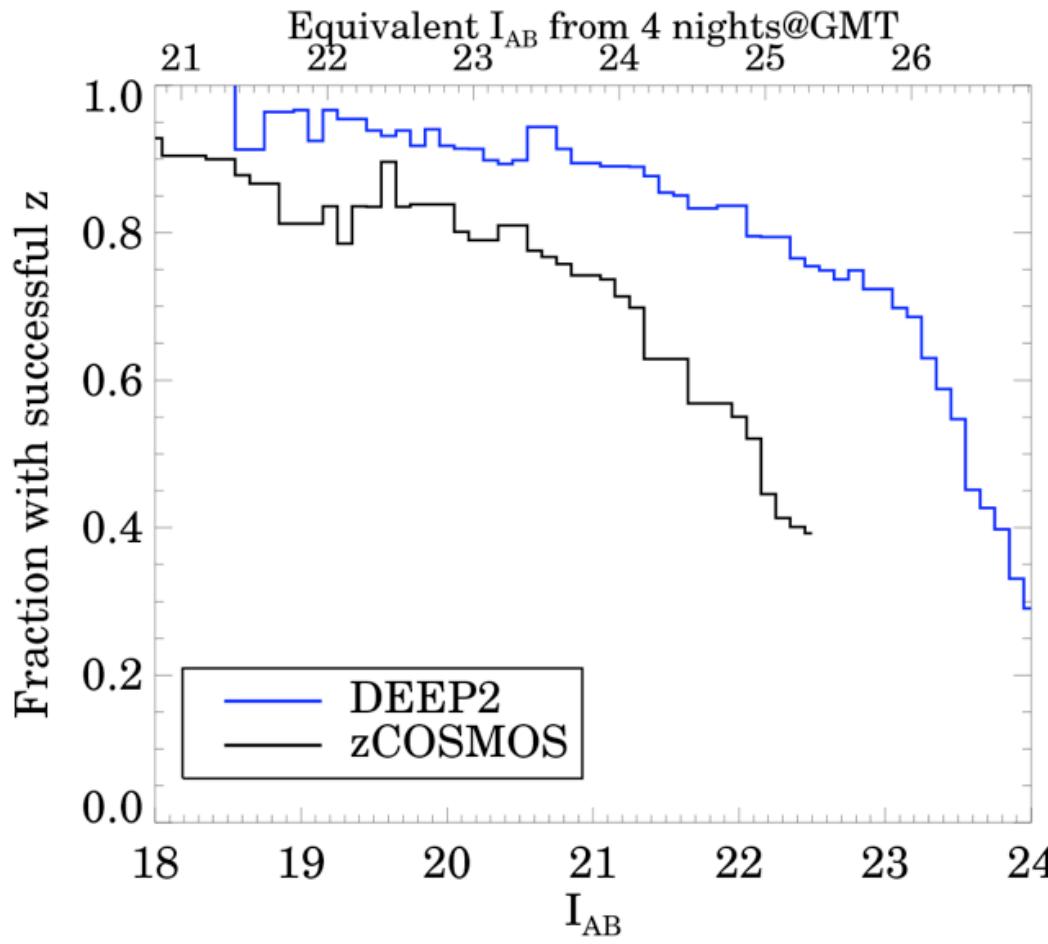


- In the next code box, change *redlimit* to a value that effectively divides red galaxies from blue ones.
- Discuss: did you get better photo-z's for this population by isolating it first?

Open issues: dealing with incompleteness in training/ calibration datasets



- In current deep spectroscopic surveys, 25-60% of targets fail to yield secure redshifts
- z success rate depends on galaxy properties
- Estimated need 99-99.9% completeness to prevent systematic errors in calibration, unless apply other methods (e.g., cross-correlations)
- Major issue for training-set techniques

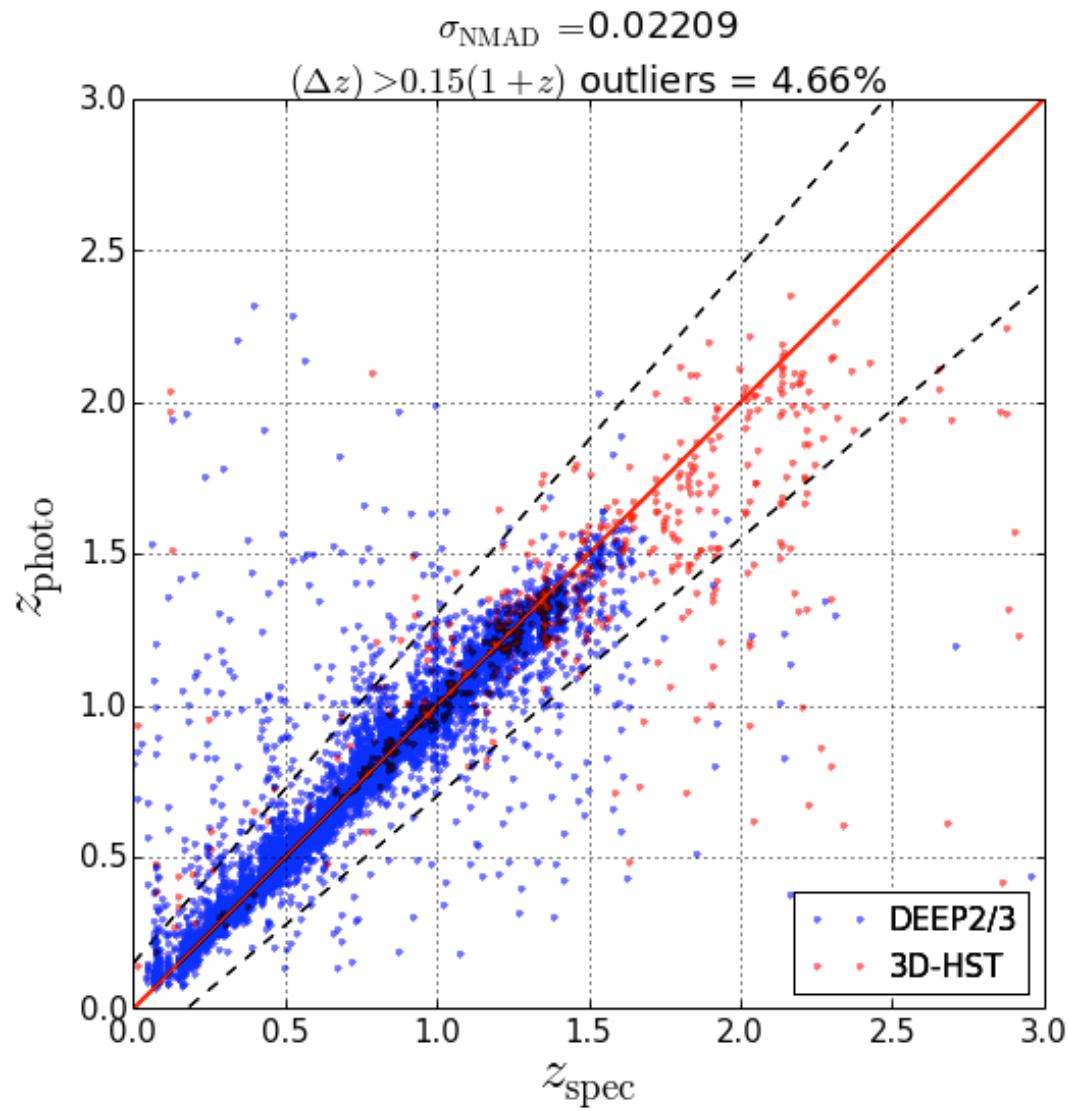


Data from DEEP2 (Newman et al. 2013) and zCOSMOS (Lilly et al. 2009)

Open issues: Robust training methods



- 1% incorrect-redshift rate is sufficient to bias photo-z's beyond tolerances
- Depending on survey, up to 5% of 'secure' redshifts are incorrect
- If can train algorithms in a manner robust to outlier/wrong redshifts, could use the broader set of less-secure spectroscopic redshifts
- ML methods that extrapolate well would also be interesting



Zhou, JN et al. 2016, in prep.

Let's explore what happens when we mis-train



- Let's try three different sorts of mis-training:
 - Training with bright objects & applying to faint (or vice versa)
 - Training with red (in observed color) objects and applying to blue
 - Training with low-z objects and applying to high-z
 - **Discuss: Which of these was worst?**
- Note: real spec-z samples have strong selection effects in redshift
- They tend to be less successful for redder and fainter galaxies as well
- Hence, our training sets are likely to suffer from all 3 problems.

Open issues: Making posteriors great again



- CANDELS code comparison: Dahlen et al. 2013
- 11 code/template combinations were tested using ~600 redshifts in GOODS-S (trained with a separate set of 600 redshifts)
- Generally χ^2 minimization, generally with some sort of prior.
- Codes with $p(z)$'s available are marked by ★

Code	Code ID	Template set	bias _{<i>z</i>} ^a	OLF ^b	σ_F^c	σ_O^d
Rainbow	A	PEGASE ^b	-0.010	0.092	0.167	0.041
GOODZ	B	CWW ^c , Kinney ^d	-0.007	0.036	0.099	0.035
EAZY ★	C	EAZY ^e +BX418 ^f	-0.009	0.051	0.114	0.044
SPOC	D	BC03 ^g	-0.030	0.147	0.197	0.073
zphot ★	E	PEGASEv2.0 ^b	-0.007	0.041	0.104	0.037
EAZY	C	EAZY ^e	-0.009	0.053	0.121	0.037
SATMC	F	BC03 ^g	-0.008	0.093	0.272	0.064
HyperZ	G	Maraston05 ^h	0.013	0.078	0.189	0.050
LePhare ★	H	BC03 ^g +Polletta07 ⁱ	-0.008	0.048	0.132	0.038
WikZ ★	I	BC03 ^g	-0.023	0.046	0.153	0.049
EAZY ★	C	EAZY ^e	-0.005	0.039	0.127	0.034
			-0.008	0.029	0.088	0.031
			-0.009	0.031	0.079	0.029

median(all)
median(5)

Open issues: Making posteriors great again



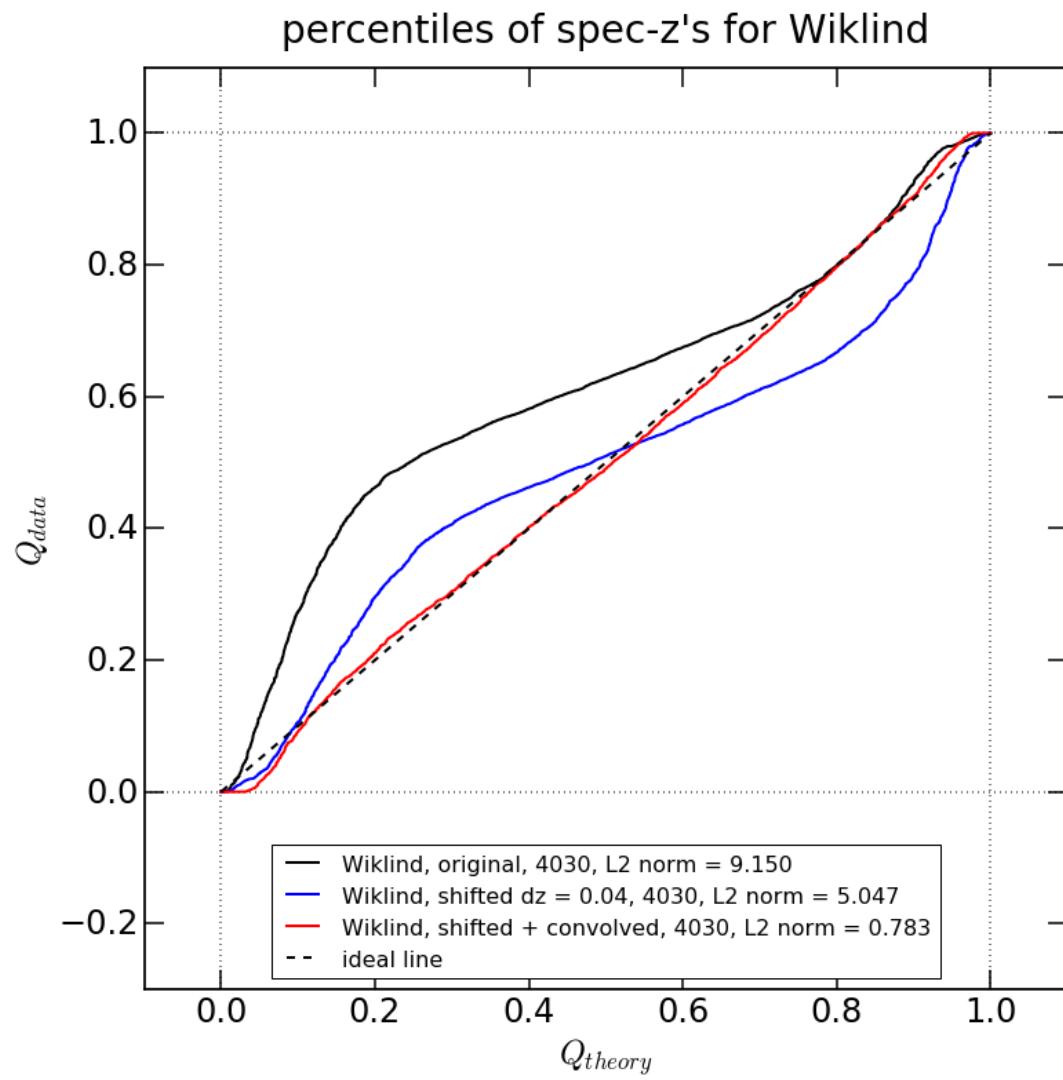
- Many analyses assume that photo-z codes are providing posterior PDFs with proper coverage (and assuming that they can add PDFs to get $N(z)$; talk to Alex Malz if you want to learn about the right way to do that...)
- Dahlen et al. 2013 tested the fraction of spectroscopic redshifts that are in the inner 68% or inner 95% of their PDFs
- Coverage is all over the place; no codes were good at both 68% and 95% points

Code conf. int:	WFC3	H-selected
	68.3%	95.4%
2A	46.1	
3B	81.6	92.8
4C	64.0	88.2
5D	2.5	4.2
6E	52.0	84.7
7C	65.0	87.3
8F	15.3	15.6
9G	16.3	44.1
11H	35.2	54.0 ^a
12I	88.7	96.7
13C	52.0	72.7

Open issues: Making posteriors great again



- LSST Dark Energy Science Collaboration is setting up controlled tests of the problem
- Meanwhile, kludge in Kodra et al. 2016: modify $p(z)$'s for CANDELS HST survey
 - Shift by constant in z direction; convolve with Gaussian kernel; and take to a power (equivalent to rescaling errors in χ^2 calculation)
- Optimize parameters by minimizing total L2 norm of deviation in Q-Q plot from expected line

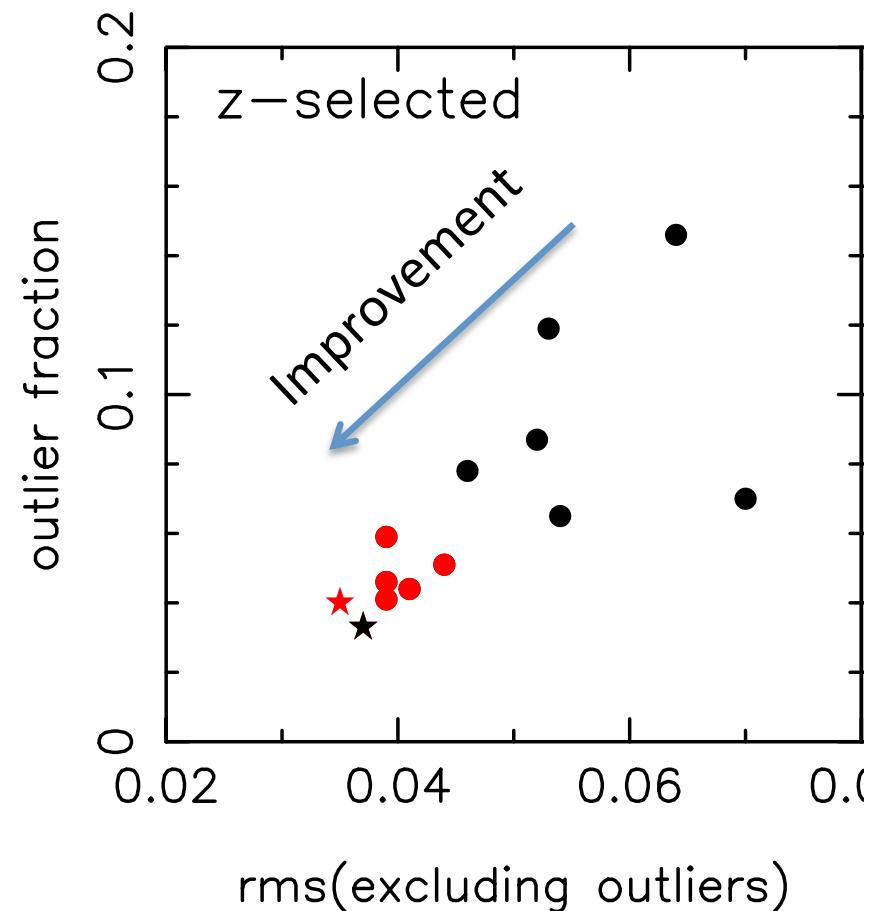


Kodra, JN et al. 2016, in prep.

Open issues: Combining PDF results from multiple codes



- Dahlen et al. found that medians of point estimates from multiple codes (\star 's) have smaller scatter (relative to spec-z) than any individual code
- All codes are run on the same data! Current codes do not make optimal use of available information...

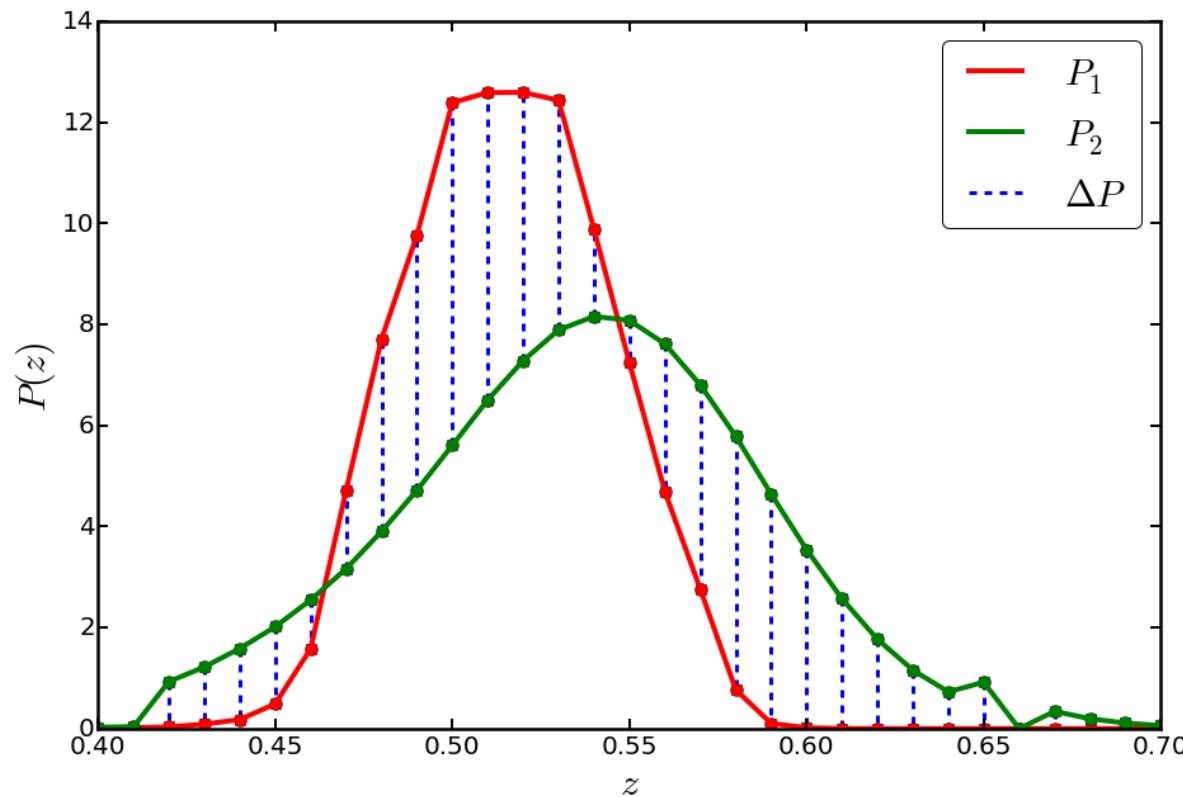


Dahlen et al. 2013

Open issues: Combining PDF results from multiple codes

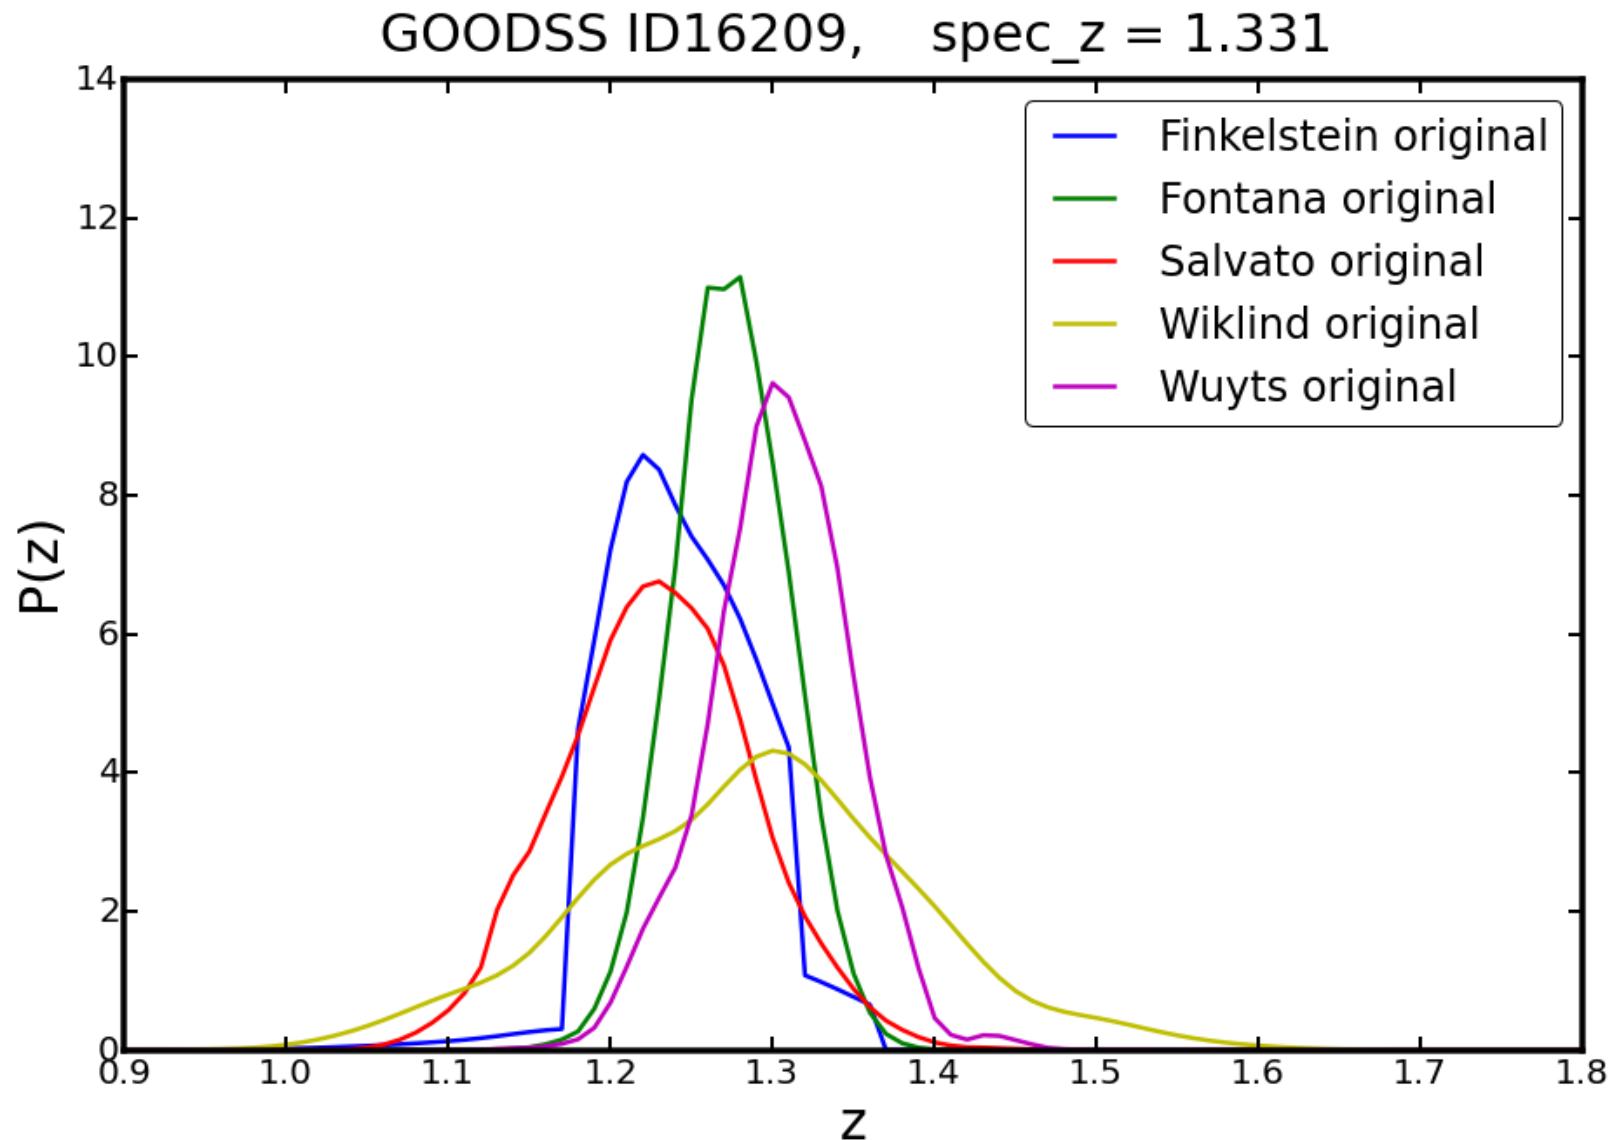


- Dahlen et al. presented a hierarchical Bayesian combination method (cf. Press & Kochanek, Lang & Hogg, etc.)
- Izbicki & Lee 2016 use weighted combinations of codes
- Kodra et al. (in prep) investigates using PDF that minimizes total Fréchet distance to remaining PDFs: analogous to median

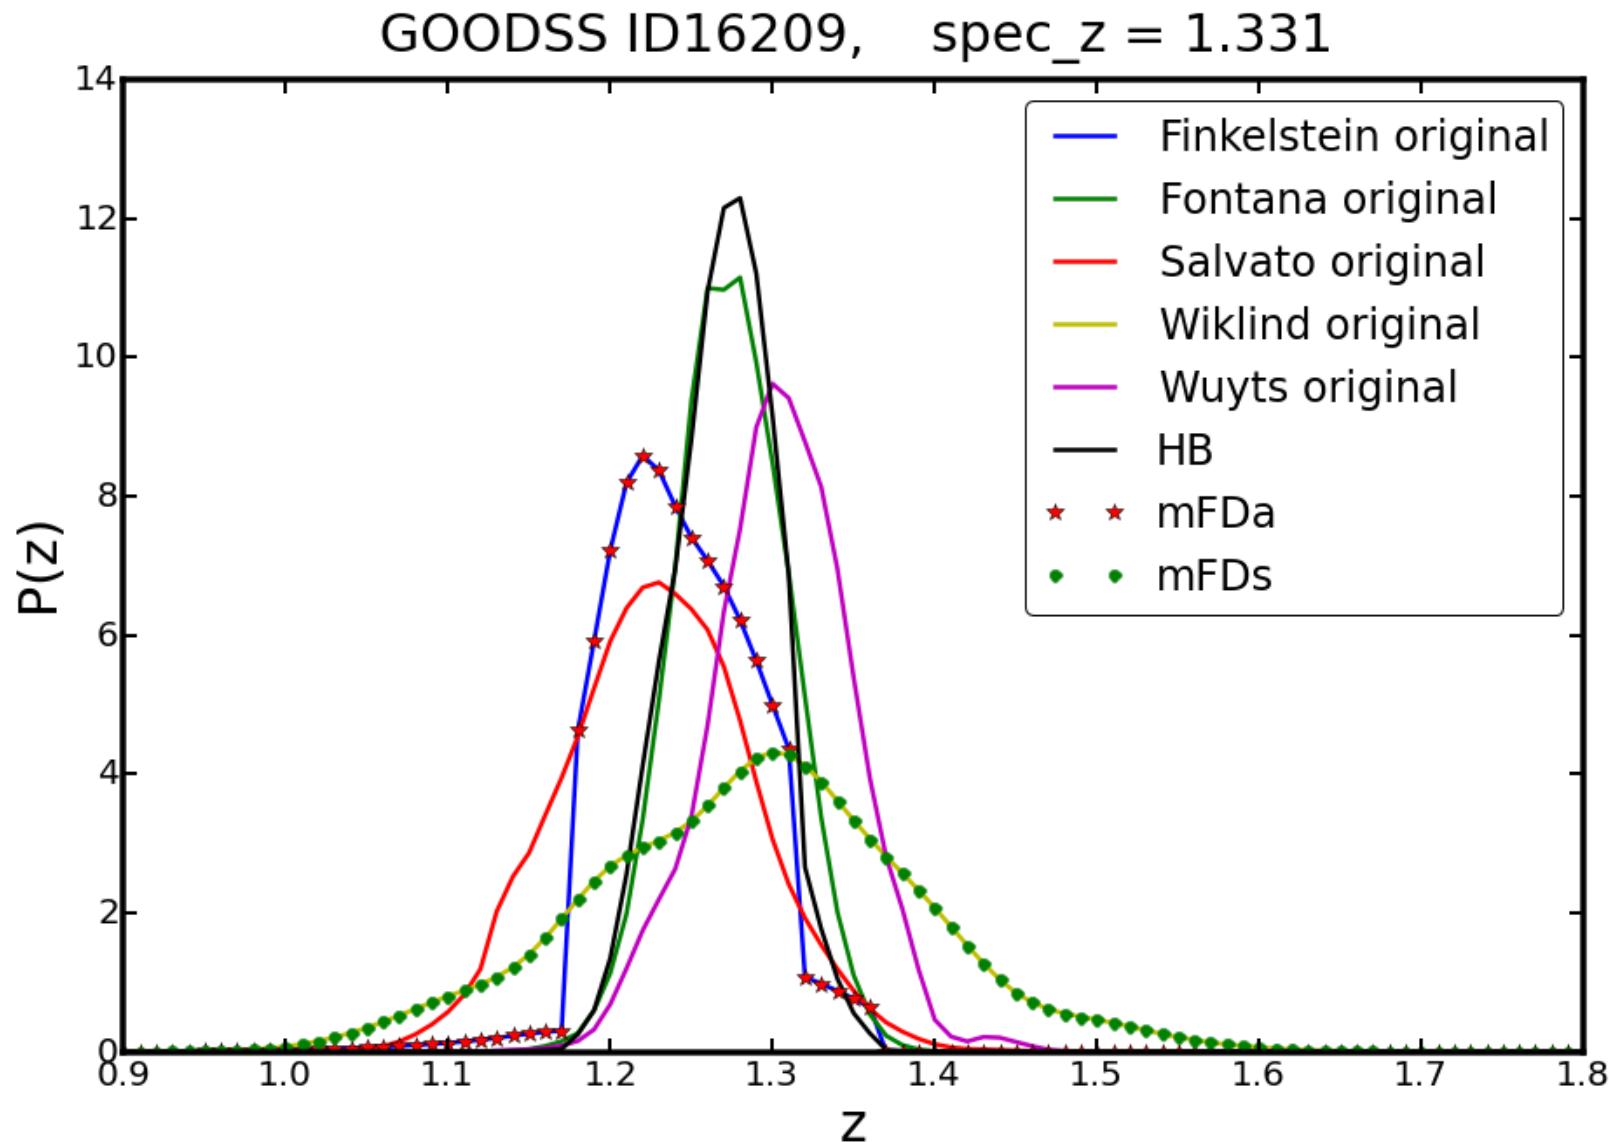


D. Kodra

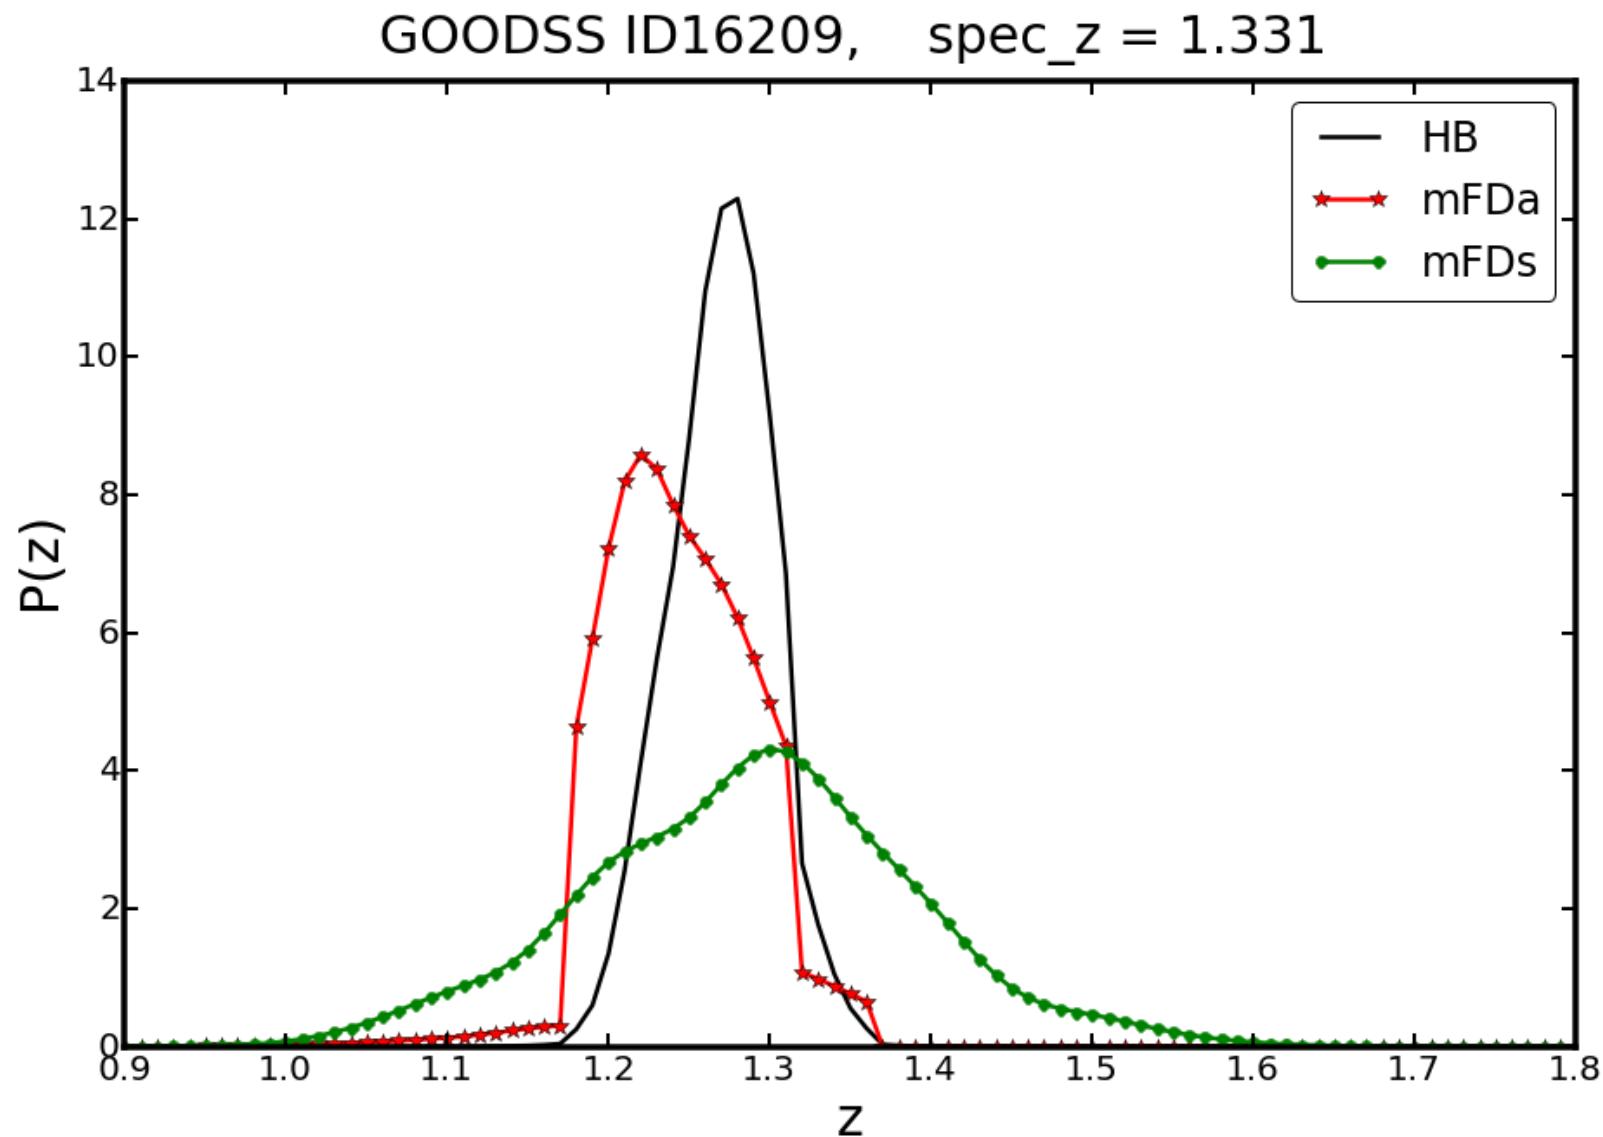
Open issues: Combining PDF results from multiple codes



Open issues: Combining PDF results from multiple codes



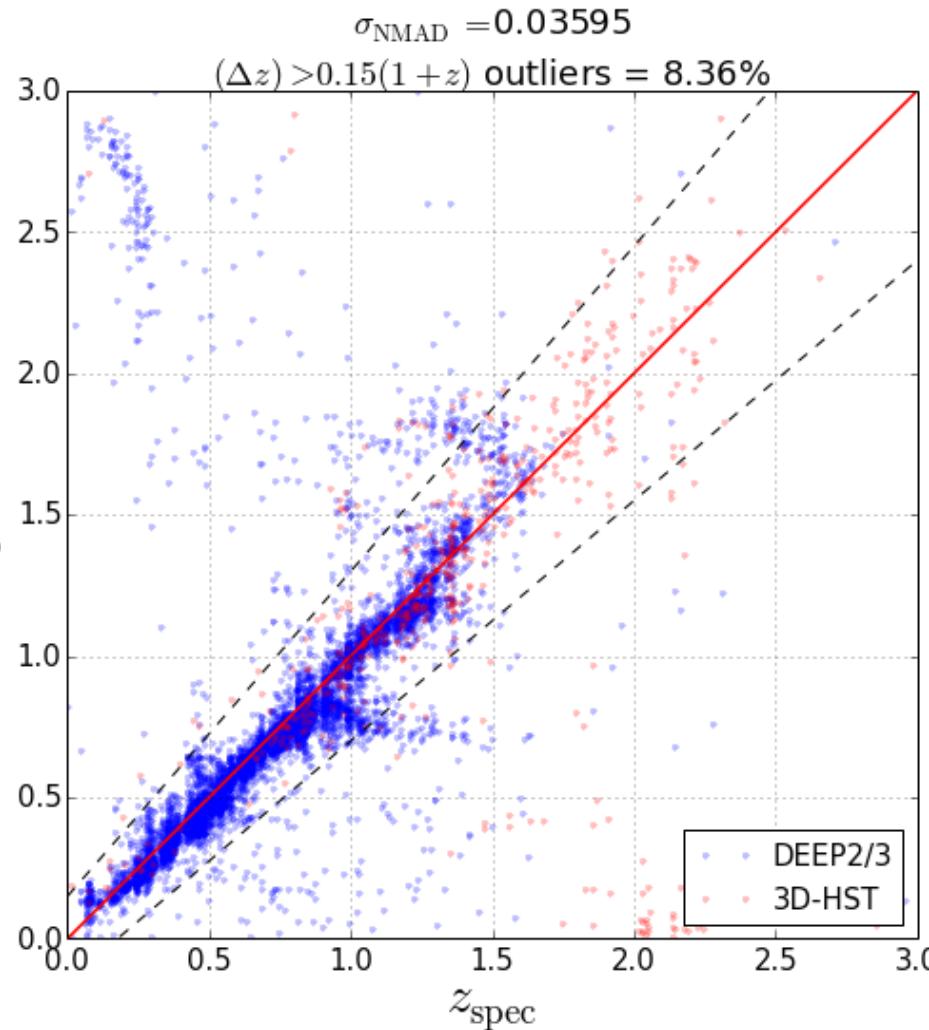
Open issues: Combining PDF results from multiple codes



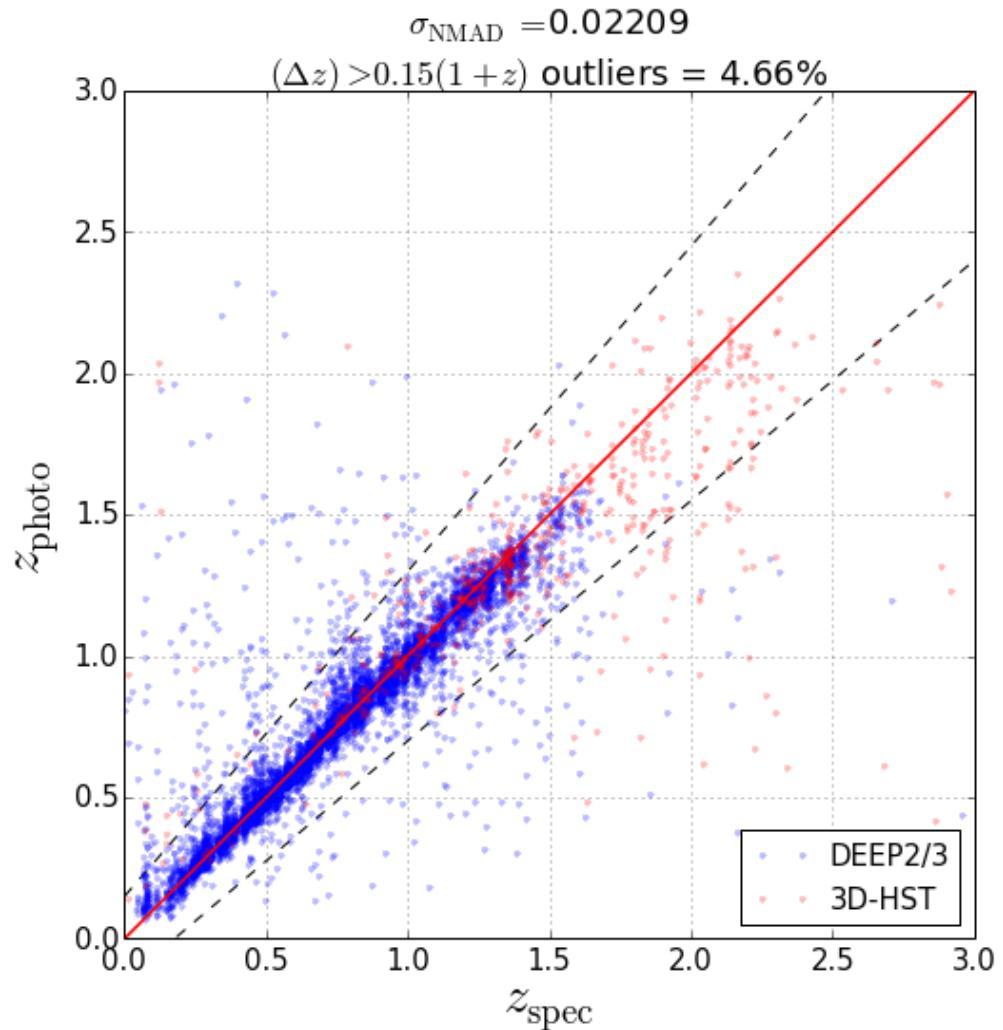
Another possible use case: template-based and training-based methods have different failure modes

- Identify potential outliers from discrepant results?

EAZY (template code, untuned)



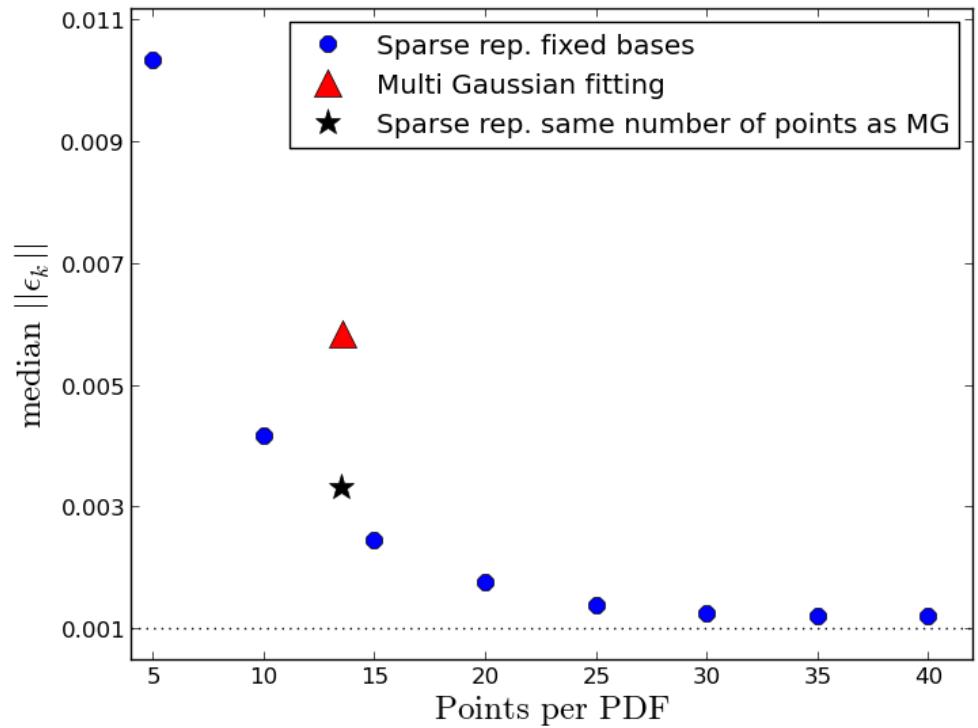
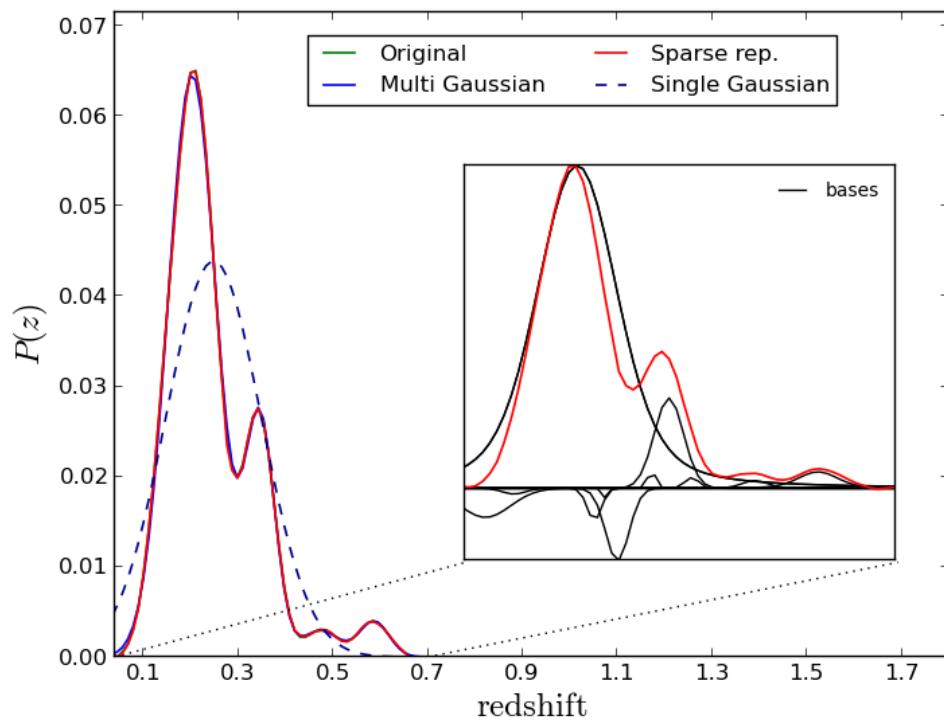
Random Forest Regression



Open issues: Storing $p(z,\alpha)$



- Carrasco-Kind & Brunner 2014 achieved strong compression of photo-z PDFs using sparse representation and well-chosen basis set
- For many LSST applications, want 2+-dimensional PDFs
- Can suitably sparse (<few hundred #s) representations be achieved?
- Are samples from PDFs OK for all science cases?

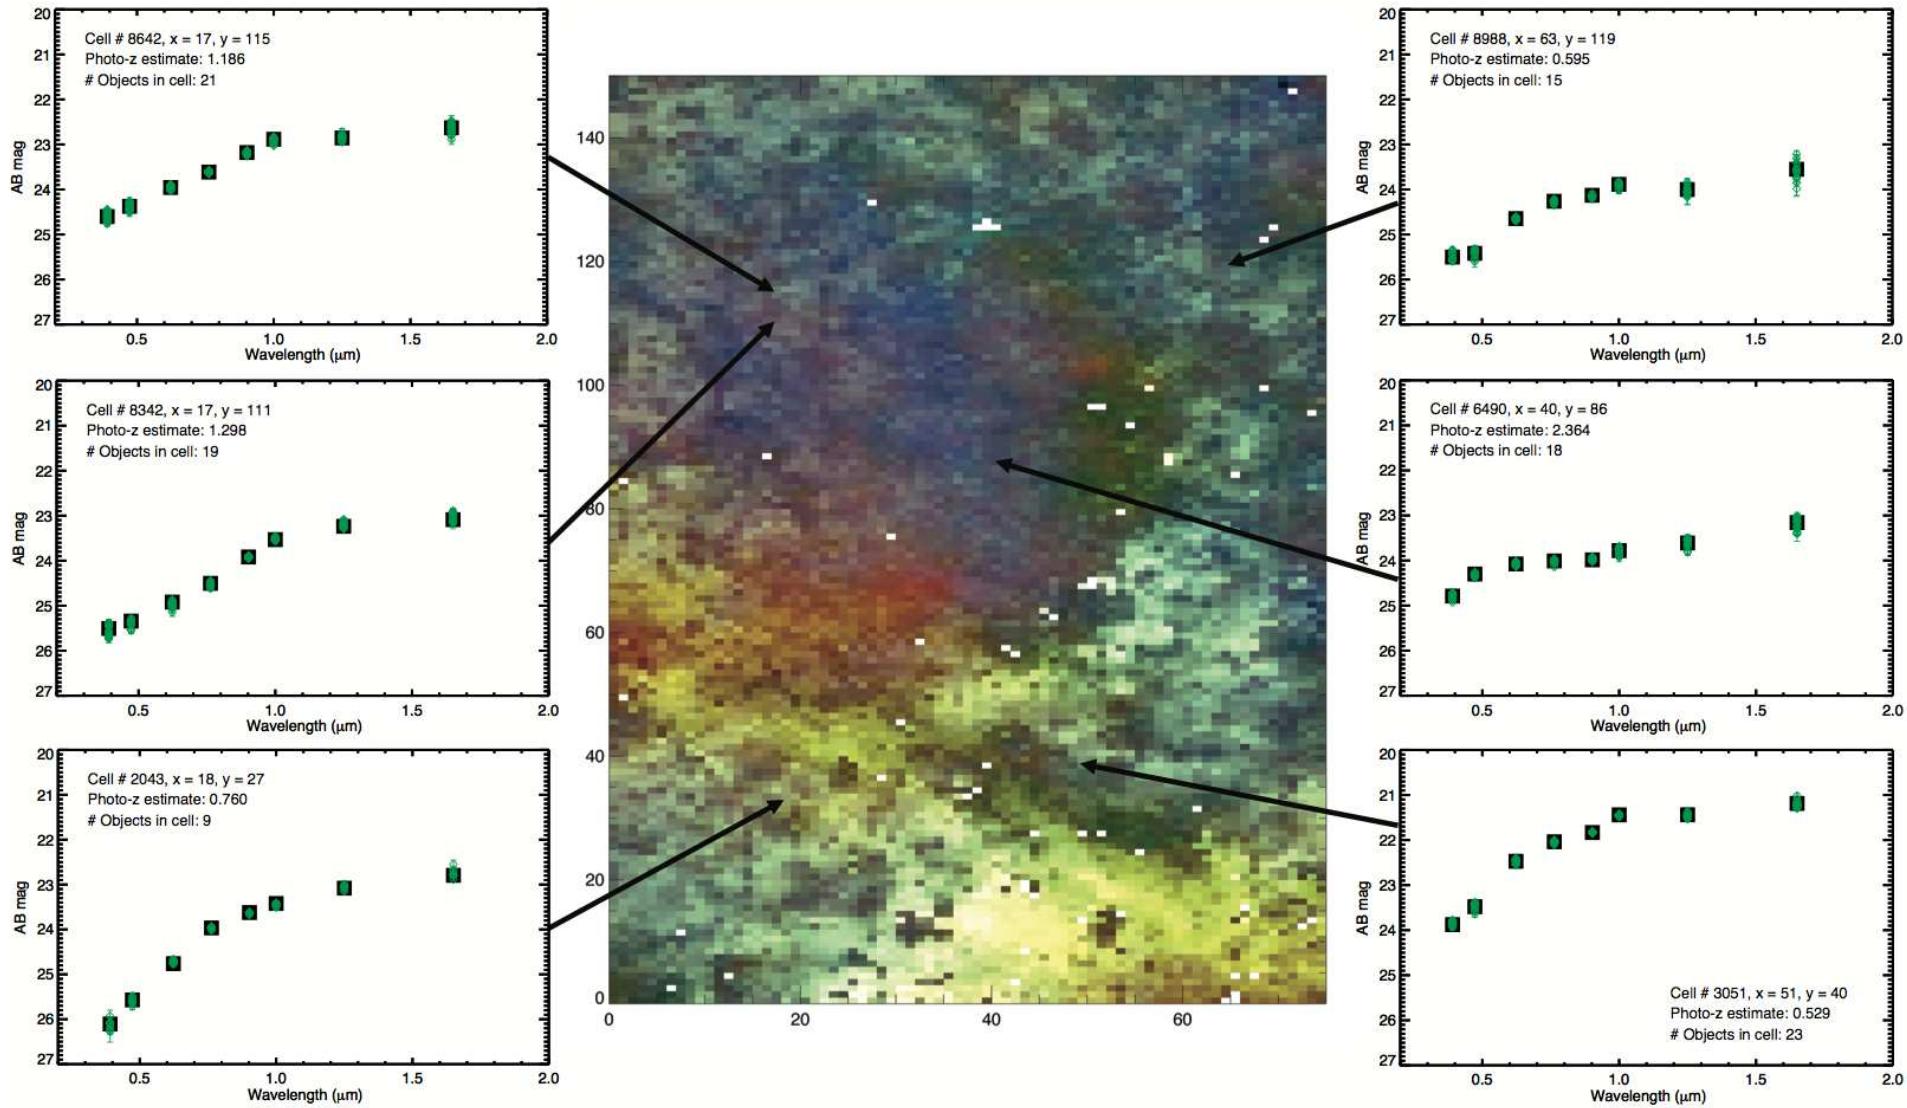


Carrasco-Kind & Brunner 2014

Open issues: Optimizing spectroscopic targeting



- Current state of the art: Masters et al. 2015
- Self-organized map of galaxy colors

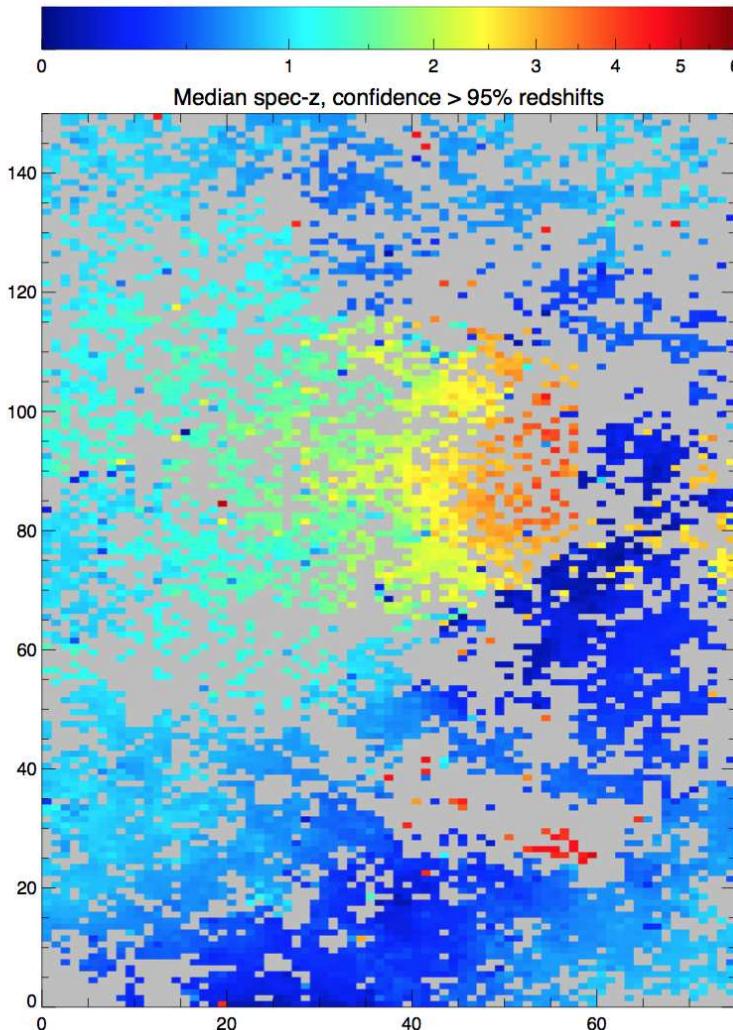
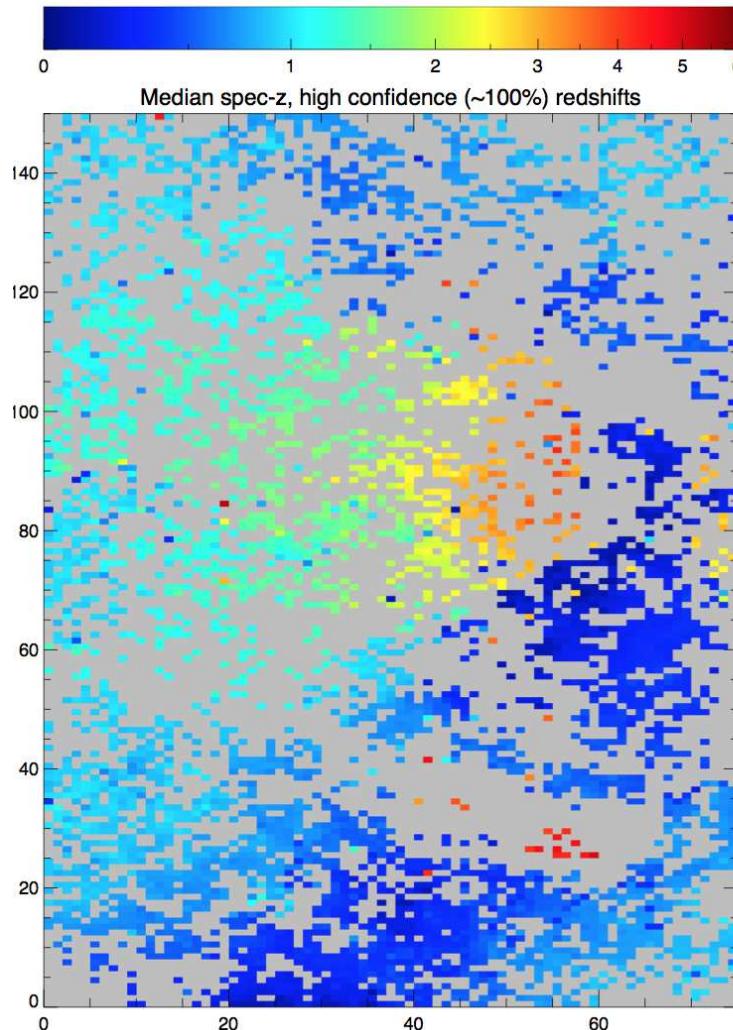


Masters et al. 2015

Open issues: Optimizing spectroscopic targeting



- Prioritize cells with few redshifts for spectroscopic follow-up
- Are there better ways to do this?



Open issues: Ideal photo-z code?

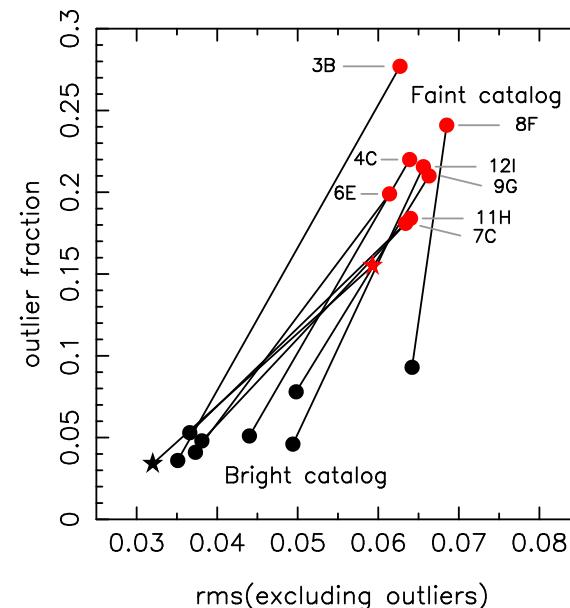
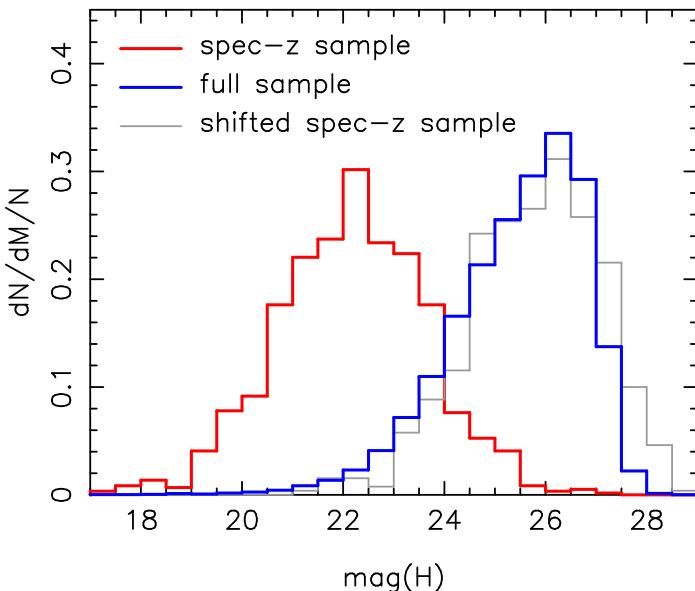


- What might an ideal LSST photo-z algorithm look like?
 - Trained with >30,000 spectra spanning range of spectra
 - Develops priors & tweaks templates via hierarchical Bayesian hyperparameters
 - Incorporates variations in effective filter wavelengths due to observational conditions: requires applying algorithm to O(1000) measurements instead of O(6)
 - Incorporates AGN classification and AGN photo-z determination: colors are not constant with time for many objects!
 - Want algorithms to be fast: create ML-based emulators for template photo-z's?
 - For bright objects, may also be useful to compare to ML techniques to identify potential outliers

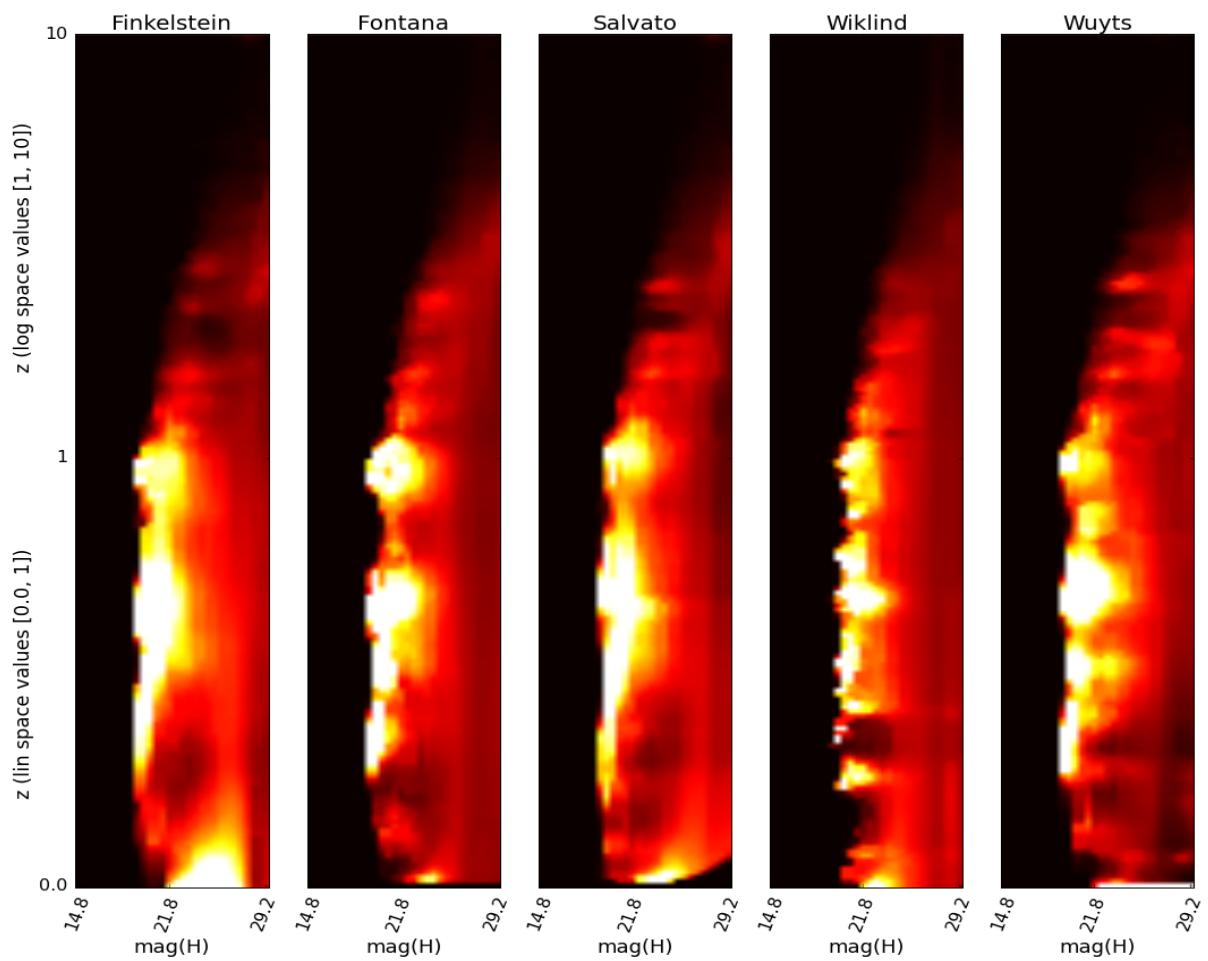
- Many tests of photo-z algorithms with deep, high-redshift dataset.

Examples:

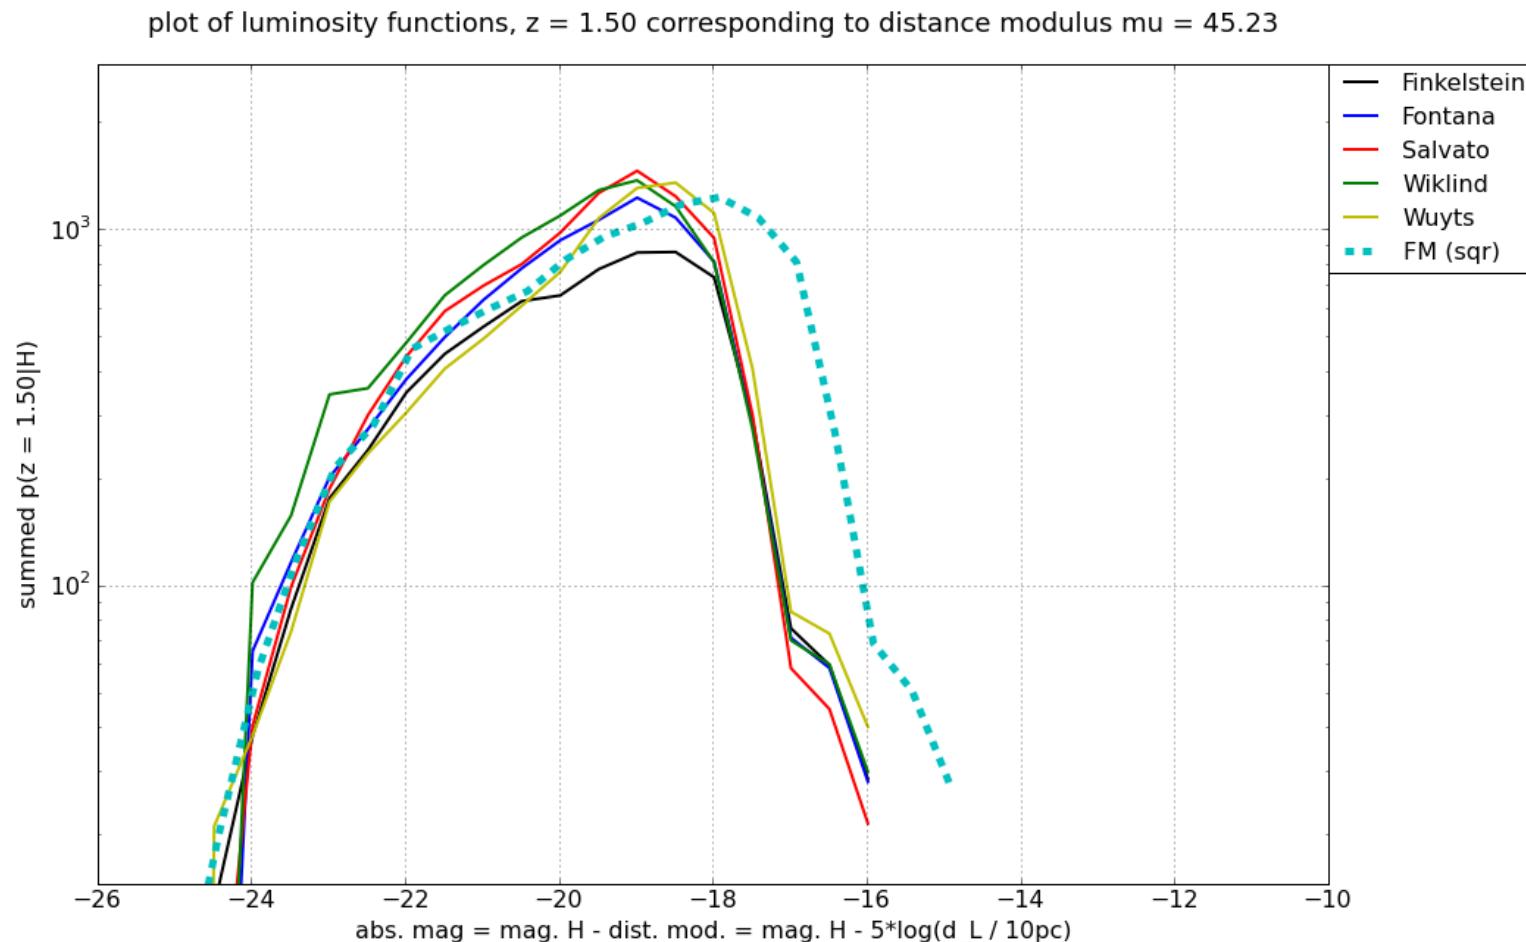
- Test photo-z performance as degrade photometry (using same test spectroscopic data)
- Dependence of errors on redshift, magnitude, & color
- Investigation of (lack of) consistency between photometric zero point shifts from different codes
- Empirical test of photo-z errors using Δz between close pairs



- Compare predictions of codes in space of $p(z \mid H)$
- Disagreement on where there are redshift spikes
- Priors have huge effect at low z (non-monotonic behavior)
- Different effective smoothings
- The performance of these codes for z_{peak} isn't all that different. . .



- This can have large (factor of few) effects on the inferred number of objects at a given redshift



Conclusions

- Training-based methods are easier to get good results from than template-based methods, but don't extrapolate well
- Key issue for LSST is inability to get complete training sets
- Other interesting issues for the next few years:
 - Training algorithms in the presence of false redshifts
 - Making sure $p(z)$'s have proper coverage
 - Combining results from multiple algorithms
 - Storing multi-dimensional PDFs compactly
 - Optimizing spectroscopic follow-up
 - Defining parameters of ideal LSST algorithm
- Current codes appear sufficient to meet LSST requirements, but are clearly suboptimal. Better photo-z's will greatly increase the value of LSST - e.g. 30% increase in Dark Energy Figure of Merit